

Name: Meet Brijwani

Roll no: 14

Batch: S11

EXPERIMENT NO : 01

Aim: Explore usage of basic linux commands and system calls for files, directory and process management.

Theory:

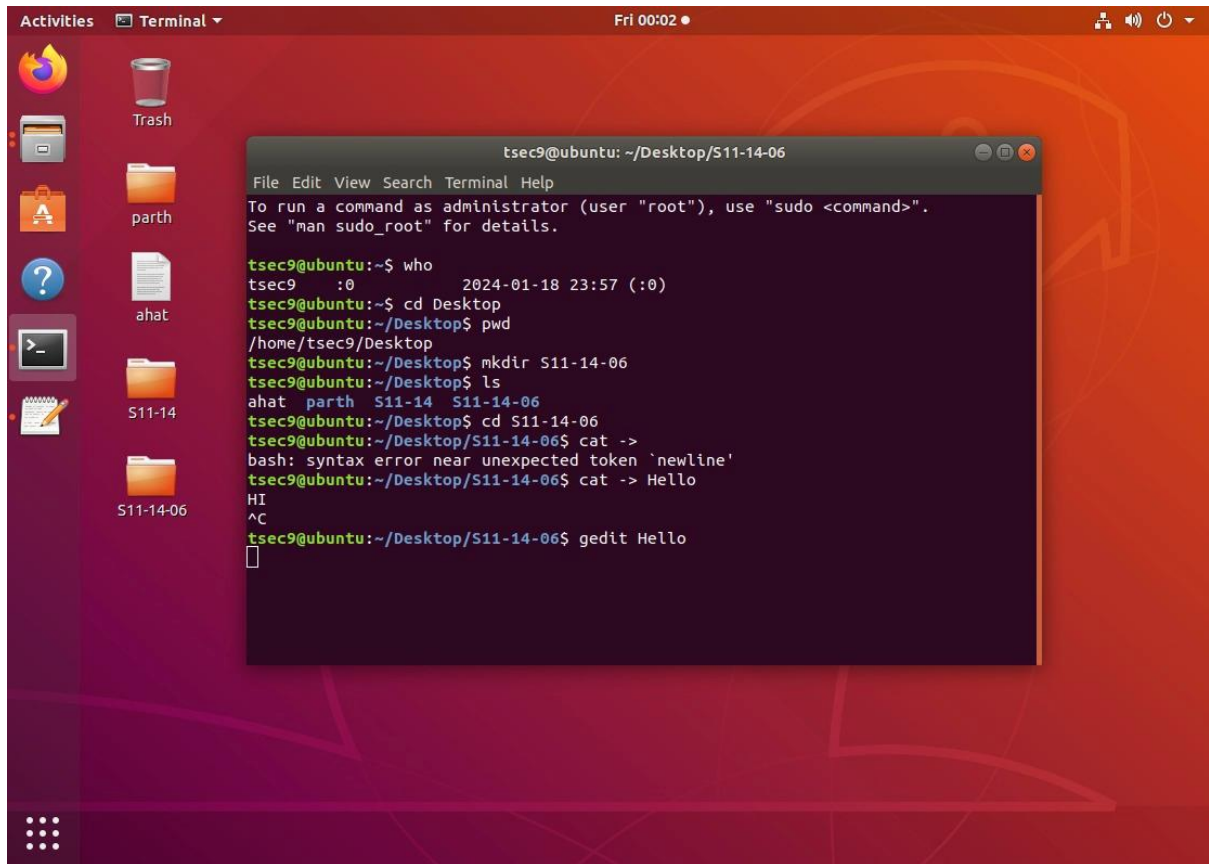
1. **who** : it is used to find out the current user who is logged into the system.
2. **pwd** : present working directory, lets you know the current directory you are in.
3. **cal** : show the calendar of the complete month.
4. **date** : It shows you the current date, along with the time, along with the day, along with the year.
5. **mkdir**: to create a new directory under any directory
6. **chdir/cd** : to change the current working directory
7. **cat** : to create the file and display the contents of the file
8. **chmod**: to change the mode of the file. There are three modes read(r), write(w) and execute(e)
9. **ls** : to list all directories and subdirectories
 - a. **ls-l** : to show the long listing information about the directory
 - b. **ls-lh** : human readable format.
 - c. **ls-ld** : shows the details of the directory content.
 - d. **ls-d*** : to show the sub directories in a directory
 - e. **ls-a** : to show hidden files
 - f. **ls-lhs** : show files in the descending order in which you have used your files.
10. **sort-r file name.txt** : sorts the list in reverse order
11. **sort-n file name.txt** : its sorts the numerical list in ascending order
12. **sort nr file name.txt** : its sorts the numerical list in reverse order
13. **sort u file name.txt** : to remove the duplicates
14. **sort m file name.txt** : Sorts the months in ascending order
15. **awk** : it is used for the user that defines text patterns that are to be searched for each line of the file.

Syntax , **awk '{print}' file name.txt**

awk '/faculty/{print}' file name.txt :

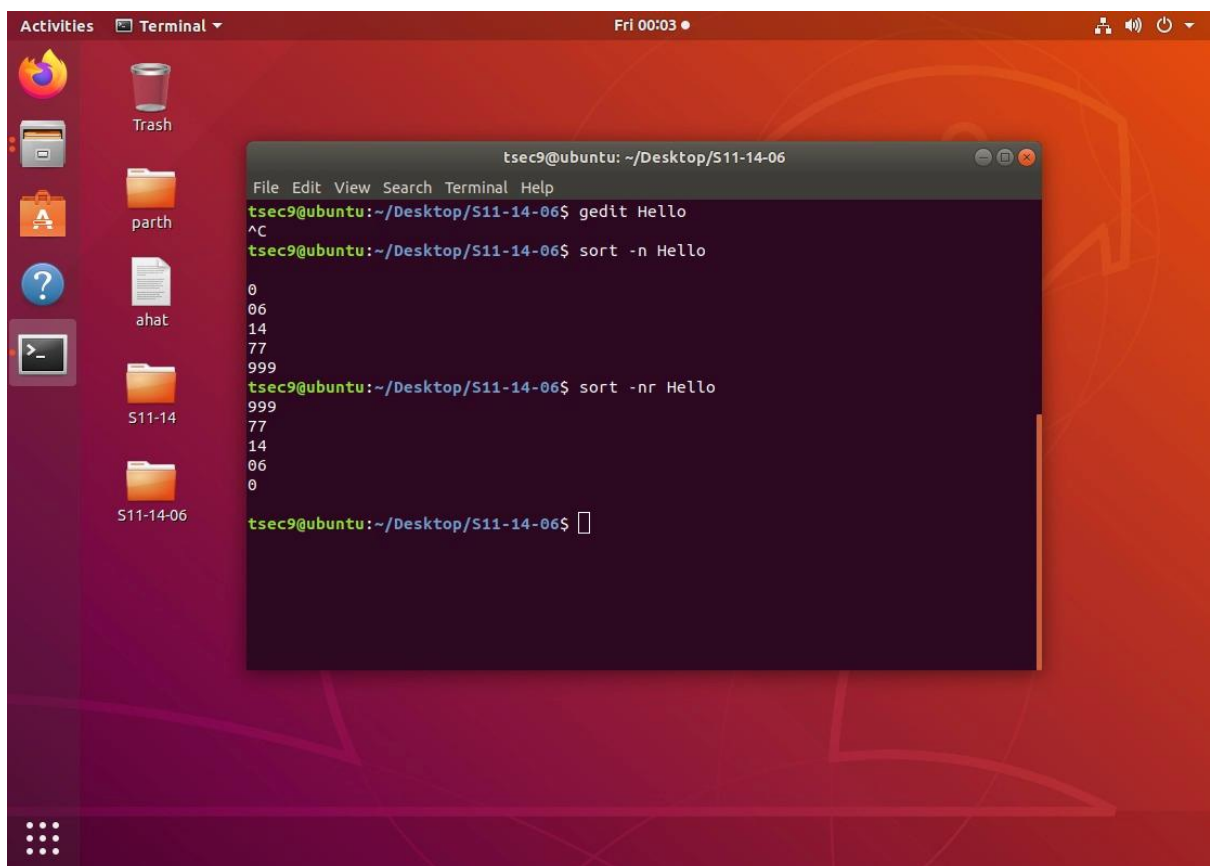
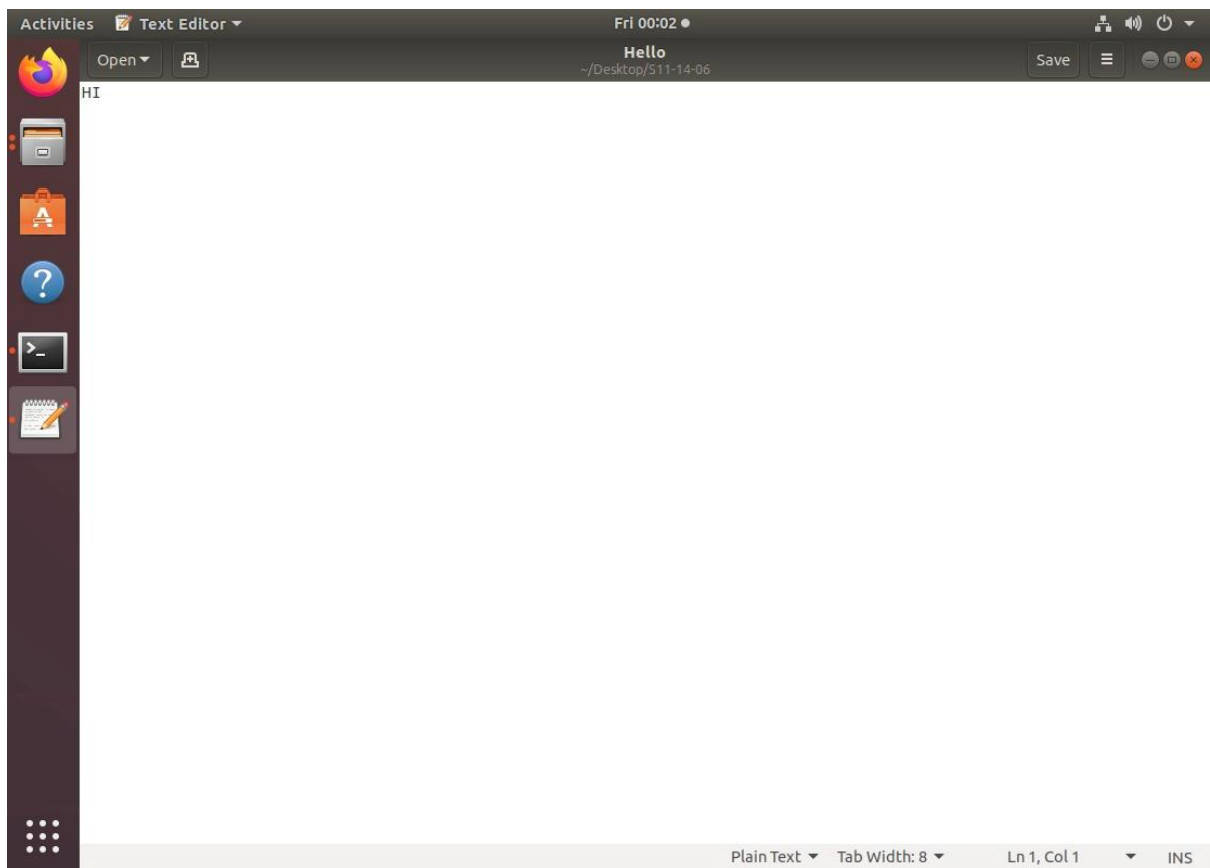
awk '{print}NR, \$0}' file name.txt :
NR - specifies the number of lines.

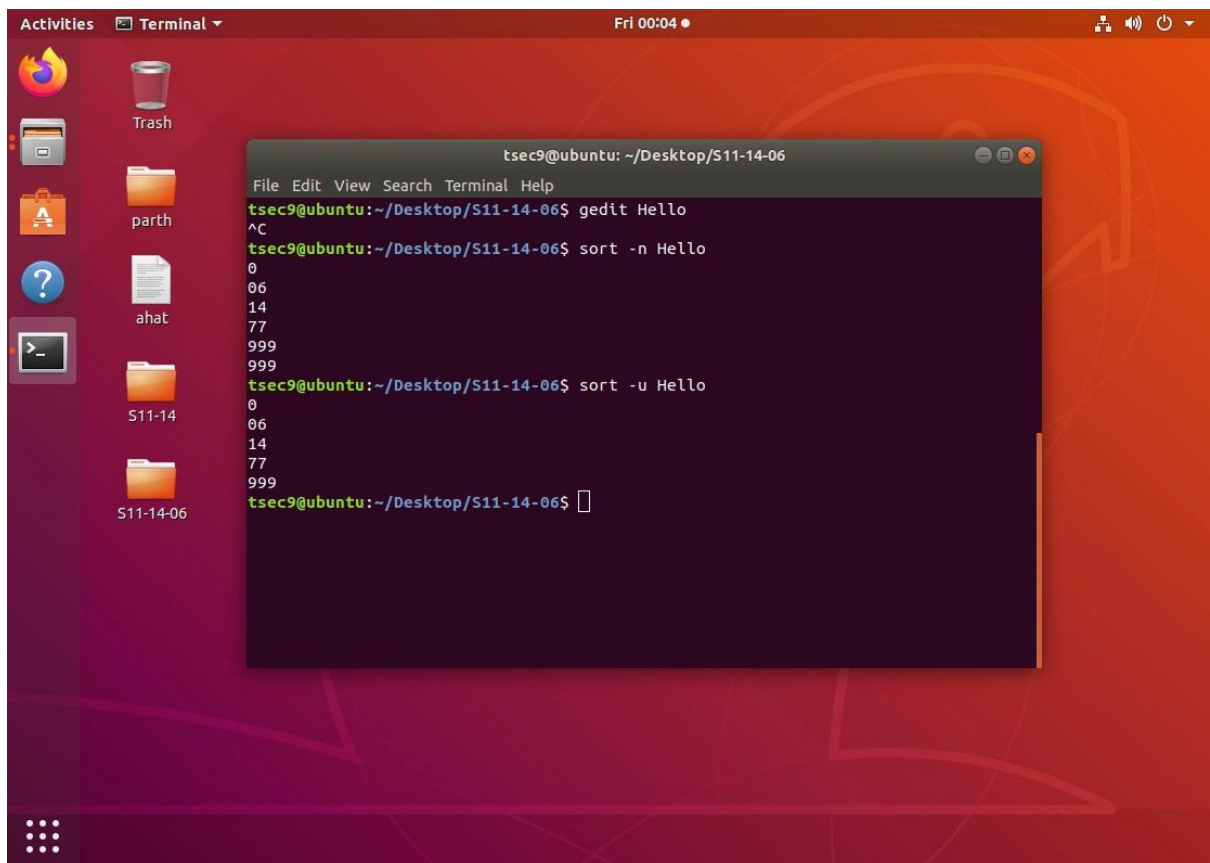
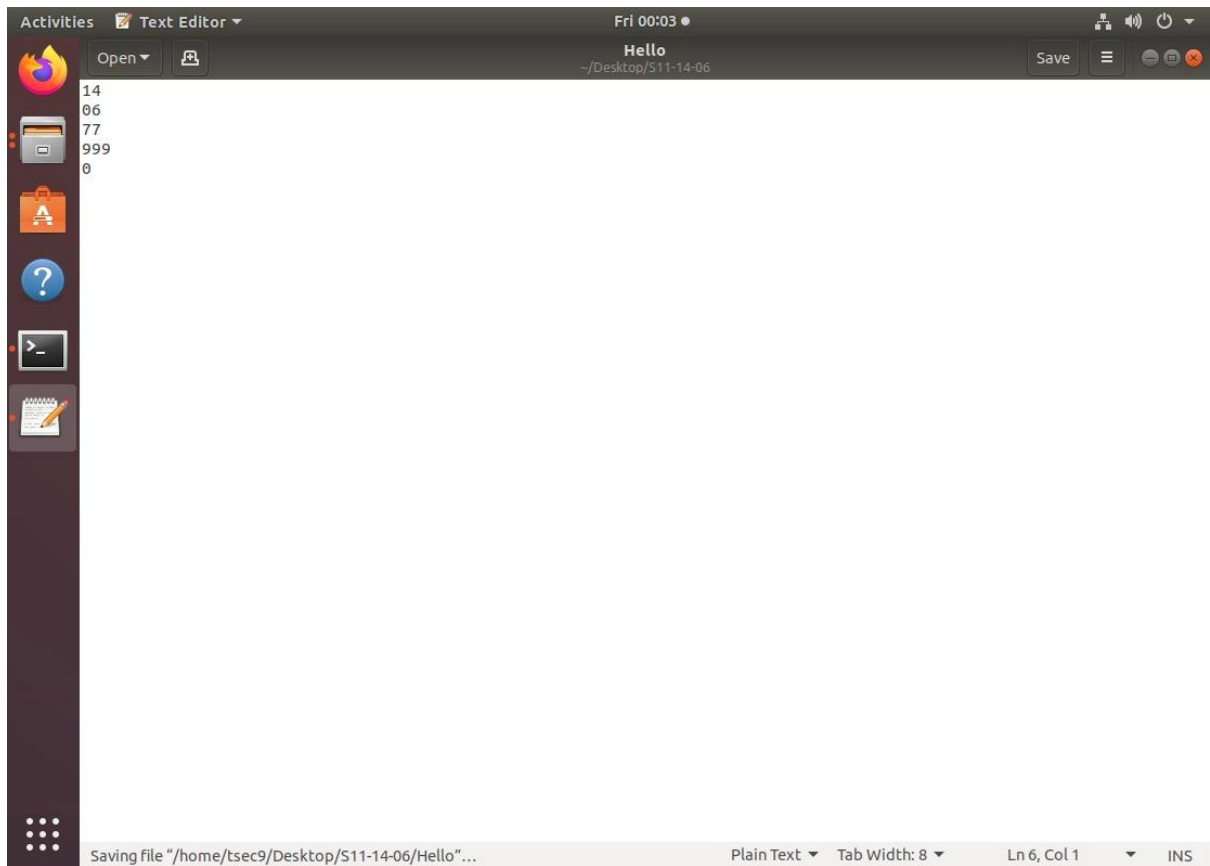
Output :



The screenshot shows an Ubuntu desktop with a red background. On the left is a dock with icons for Firefox, Trash, parth, ahat, a terminal, and two folders named S11-14 and S11-14-06. The top bar shows 'Activities', 'Terminal', and the time 'Fri 00:02'. A terminal window titled 'tsec9@ubuntu: ~/Desktop/S11-14-06' is open, displaying the following commands and output:

```
tsec9@ubuntu:~$ who
tsec9    :0                2024-01-18 23:57 (:0)
tsec9@ubuntu:~$ cd Desktop
tsec9@ubuntu:~/Desktop$ pwd
/home/tsec9/Desktop
tsec9@ubuntu:~/Desktop$ mkdir S11-14-06
tsec9@ubuntu:~/Desktop$ ls
ahat  parth  S11-14  S11-14-06
tsec9@ubuntu:~/Desktop$ cd S11-14-06
tsec9@ubuntu:~/Desktop/S11-14-06$ cat ->
bash: syntax error near unexpected token `newline'
tsec9@ubuntu:~/Desktop/S11-14-06$ cat -> Hello
HI
^C
tsec9@ubuntu:~/Desktop/S11-14-06$ gedit Hello
```





Name: Meet Brijwani

Roll no: 14

Batch: S11

Experiment 2: Linux shell script

Aim:

Write shell scripts to do the following:

- a. Display OS version, release number, kernel version
- b. Display top 10 processes in descending order
- c. Display processes with highest memory usage.
- d. Display current logged in user and log name.

Display current shell, home directory, operating system type, current path setting, current working directory.

Theory:

A]

i) `uname - r {username}`

To show the OS version, release no. OS release version

ii) `uname - v`

Shows Kernel version (last when launched)

iii) `uname - a`

Shows all the details including OS used, PC and other details.

`us - c + u - r → us - a`

B]

i) `ps-aux (i) sort : nk +41 tail`

Sorts 10 processes in descending order. It shows logged in users, modes of memory & storage.

OR

ii) `ps-aux | tail`

Same as above

OR

iii) `ps-aux | sort -nl+4 | tail - n15` (shows top 15 processes)
same with 15

C]

i) `sudo apt install htop`

shows date of processes with highest memory

ii) `ps - eo pid,ppid,cmd,%mem,%cpu --sort=-%mem | head`
same as above same as above

D]

i) `ps-p$$`

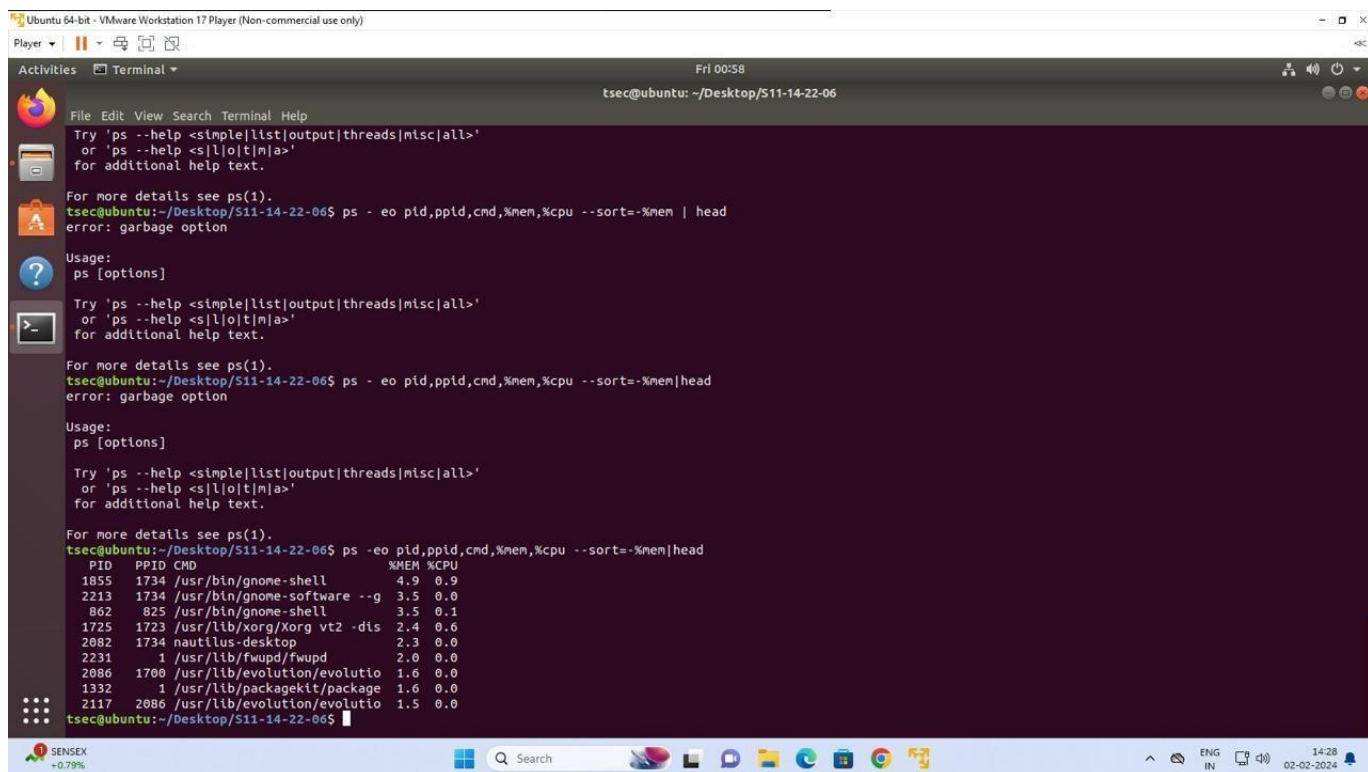
Shows process ID and current shell

ii) `echo display home directory`

`echo $ home`

To display home directory

Output:



The screenshot shows a terminal window titled "tsec@ubuntu: ~/Desktop/S11-14-22-06". The terminal displays the output of the command `ps -eo pid,ppid,cmd,%mem,%cpu --sort=-%mem | head`. The output is a table with columns: PID, PPID, CMD, %MEM, and %CPU. The processes are sorted by memory usage in descending order.

PID	PPID	CMD	%MEM	%CPU
1855	1734	/usr/bin/gnome-shell	4.9	0.9
2213	1734	/usr/bin/gnome-software --g	3.5	0.0
862	825	/usr/bin/gnome-shell	3.5	0.1
1725	1723	/usr/lib/Xorg/Xorg vt2 -dis	2.4	0.6
2082	1734	nautilus-desktop	2.3	0.0
2231	1	/usr/lib/fwupd/fwupd	2.0	0.0
2086	1700	/usr/lib/evolution/evolutio	1.6	0.0
1332	1	/usr/lib/packagekit/package	1.6	0.0
2117	2086	/usr/lib/evolution/evolutio	1.5	0.0


```
Ubuntu 64-bit - VMware Workstation 17 Player (Non-commercial use only)
Player
Fri 00:36
tsec@ubuntu: ~/Desktop/S11-14-22-06

File Edit View Search Terminal Help
tsec@ubuntu:~/Desktop/S11-14-22-06$ clear
tsec@ubuntu:~/Desktop/S11-14-22-06$ cat -> EnvironmentVar.sh
^C
tsec@ubuntu:~/Desktop/S11-14-22-06$ ls
EnvironmentVar.sh
tsec@ubuntu:~/Desktop/S11-14-22-06$ gedit EnvironmentVar.sh
tsec@ubuntu:~/Desktop/S11-14-22-06$ sh EnvironmentVar.sh
Path for the system is: /usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/games:/usr/local/games:/snap/bin
Display User ID as:
Display the parent ProCessID: 2381
Display the username:
Display the current directory: /home/tsec/Desktop/S11-14-22-06
Disk Usage
Free/Total disk:1.9G/1.9G
Top 10 processes in the system sorted by memory
MemTotal: 3994712 kB
MemFree: 748436 kB
MemAvailable: 2285584 kB
Buffers: 179988 kB
Cached: 1521068 kB
SwapCached: 0 kB
Active: 2059504 kB
Inactive: 556276 kB
Active(anon): 772768 kB
Inactive(anon): 162152 kB
Active(file): 1286736 kB
Inactive(file): 394124 kB
Unevictable: 16 kB
Mlocked: 16 kB
SwapTotal: 969960 kB
SwapFree: 968924 kB
Dirty: 72 kB
Writeback: 0 kB
AnonPages: 914460 kB
Mapped: 241224 kB
Shmem: 20352 kB
KReclaimable: 122292 kB
Slab: 209544 kB
SReclaimable: 122292 kB
SUnreclaim: 87252 kB
```

```
Ubuntu 64-bit - VMware Workstation 17 Player (Non-commercial use only)
Player
Fri 00:36
tsec@ubuntu: ~/Desktop/S11-14-22-06

File Edit View Search Terminal Help
WritebackTmp: 0 kB
CommitLimit: 2967316 kB
Committed_AS: 4885052 kB
VmallocTotal: 34359738367 kB
VmallocUsed: 27652 kB
VmallocChunk: 0 kB
Percpu: 53760 kB
HardwareCorrupted: 0 kB
AnonHugePages: 0 kB
ShmemHugePages: 0 kB
ShmemPmdMmapped: 0 kB
FileHugePages: 0 kB
FilePmdMmapped: 0 kB
CmaTotal: 0 kB
CmaFree: 0 kB
HugePages_Total: 0
HugePages_Free: 0
HugePages_Rsvd: 0
HugePages_Surp: 0
HugePagesize: 2048 kB
Hugetlb: 0 kB
DirectMap4k: 259904 kB
DirectMap2M: 3934208 kB
DirectMap1G: 2097152 kB
The OS version is:
NAME="Ubuntu"
VERSION="18.04.6 LTS (Bionic Beaver)"
ID=ubuntu
ID_LIKE=debian
PRETTY_NAME="Ubuntu 18.04.6 LTS"
VERSION_ID="18.04"
HOME_URL="https://www.ubuntu.com/"
SUPPORT_URL="https://help.ubuntu.com/"
BUG_REPORT_URL="https://bugs.launchpad.net/ubuntu/"
PRIVACY_POLICY_URL="https://www.ubuntu.com/legal/terms-and-policies/privacy-policy"
VERSION_CODENAME=bionic
UBUNTU_CODENAME=bionic
The kernel version is:
Linux version 5.4.0-150-generic (builddg@bos03-amd64-012) (gcc version 7.5.0 (Ubuntu 7.5.0-3ubuntu1-18.04)) #167-18.04.1-Ubuntu SMP Wed May 24 00:51:42 UTC 2023
tsec@ubuntu:~/Desktop/S11-14-22-06$ gedit EnvironmentVar.sh
```



```
Ubuntu 64-bit - VMware Workstation 17 Player (Non-commercial use only)
Player
Activities Terminal
Fri 00:50
tsec@ubuntu: ~/Desktop/S11-14-22-06

File Edit View Search Terminal Help
ps [options]

Try 'ps --help <simple|list|output|threads|misc|all>'
or 'ps --help <s|l|o|t|m|a>'
for additional help text.

For more details see ps(1).
tsec@ubuntu:~/Desktop/S11-14-22-06$ ps -eo pid, ppid, cmd, %mem, %cpu -sort ==%mem | head
error: improper list

Usage:
ps [options]

Try 'ps --help <simple|list|output|threads|misc|all>'
or 'ps --help <s|l|o|t|m|a>'
for additional help text.

For more details see ps(1).
tsec@ubuntu:~/Desktop/S11-14-22-06$ uname -v
#167-18.04.1-Ubuntu SMP Wed May 24 00:51:42 UTC 2023
tsec@ubuntu:~/Desktop/S11-14-22-06$ uname -a
Linux ubuntu 5.4.0-150-generic #167-18.04.1-Ubuntu SMP Wed May 24 00:51:42 UTC 2023 x86_64 x86_64 x86_64 GNU/Linux
tsec@ubuntu:~/Desktop/S11-14-22-06$ name -r(username)

Command 'name' not found, did you mean:

command 'name' from snap name (name0261)
command 'uname' from deb coreutils
command 'name' from deb name
command 'nvme' from deb nvme-cli
command 'lane' from deb lane
command 'nanel' from deb util-linux
command 'nam' from deb nam
command 'nama' from deb nama
command 'named' from deb bind9

See 'snap info <snapname>' for additional versions.
tsec@ubuntu:~/Desktop/S11-14-22-06$
```

```
Ubuntu 64-bit - VMware Workstation 17 Player (Non-commercial use only)
Player
Activities Terminal
Fri 00:51
tsec@ubuntu: ~/Desktop/S11-14-22-06

File Edit View Search Terminal Help
Try 'ps --help <simple|list|output|threads|misc|all>'
or 'ps --help <s|l|o|t|m|a>'
for additional help text.

For more details see ps(1).
tsec@ubuntu:~/Desktop/S11-14-22-06$ ps -eo pid, ppid, cmd, %mem, %cpu -sort ==%mem | head
error: improper list

Usage:
ps [options]

Try 'ps --help <simple|list|output|threads|misc|all>'
or 'ps --help <s|l|o|t|m|a>'
for additional help text.

For more details see ps(1).
tsec@ubuntu:~/Desktop/S11-14-22-06$ uname -v
#167-18.04.1-Ubuntu SMP Wed May 24 00:51:42 UTC 2023
tsec@ubuntu:~/Desktop/S11-14-22-06$ uname -a
Linux ubuntu 5.4.0-150-generic #167-18.04.1-Ubuntu SMP Wed May 24 00:51:42 UTC 2023 x86_64 x86_64 x86_64 GNU/Linux
tsec@ubuntu:~/Desktop/S11-14-22-06$ name -r(username)

Command 'name' not found, did you mean:

command 'name' from snap name (name0261)
command 'uname' from deb coreutils
command 'name' from deb name
command 'nvme' from deb nvme-cli
command 'lane' from deb lane
command 'nanel' from deb util-linux
command 'nam' from deb nam
command 'nama' from deb nama
command 'named' from deb bind9

See 'snap info <snapname>' for additional versions.
tsec@ubuntu:~/Desktop/S11-14-22-06$ uname -r
5.4.0-150-generic
tsec@ubuntu:~/Desktop/S11-14-22-06$
```

```
Ubuntu 64-bit - VMware Workstation 17 Player (Non-commercial use only)
Player
Activities Terminal
Fri 00:53
tsec@ubuntu: ~/Desktop/S11-14-22-06

File Edit View Search Terminal Help
tsec@ubuntu:~/Desktop/S11-14-22-06$ ps -eo pid,ppid,cnd,%mem,%cpu -sort ==%mem | head
error: improper list

Usage:
ps [options]

Try 'ps --help <simple|list|output|threads|misc|all>'
or 'ps --help <s|l|o|t|m|a>'
for additional help text.

For more details see ps(1).
tsec@ubuntu:~/Desktop/S11-14-22-06$ uname -v
#167-18.04.1-Ubuntu SMP Wed May 24 00:51:42 UTC 2023
tsec@ubuntu:~/Desktop/S11-14-22-06$ uname -a
Linux ubuntu 5.4.0-150-generic #167-18.04.1-Ubuntu SMP Wed May 24 00:51:42 UTC 2023 x86_64 x86_64 x86_64 GNU/Linux
tsec@ubuntu:~/Desktop/S11-14-22-06$ name -r{username}

Command 'name' not found, did you mean:

command 'name' from snap name (name0261)
command 'uname' from deb coreutils
command 'name' from deb name
command 'nvme' from deb nvme-cli
command 'lane' from deb lane
command 'names' from deb util-linux
command 'nam' from deb nam
command 'nana' from deb nana
command 'named' from deb bind9

See 'snap info <snapname>' for additional versions.
tsec@ubuntu:~/Desktop/S11-14-22-06$ uname -r
5.4.0-150-generic
tsec@ubuntu:~/Desktop/S11-14-22-06$ echo $ home
$ home
tsec@ubuntu:~/Desktop/S11-14-22-06$ echo $
$
tsec@ubuntu:~/Desktop/S11-14-22-06$
tsec@ubuntu:~/Desktop/S11-14-22-06$
```

```
Ubuntu 64-bit - VMware Workstation 17 Player (Non-commercial use only)
Player
Activities Terminal
Fri 00:58
tsec@ubuntu: ~/Desktop/S11-14-22-06

File Edit View Search Terminal Help
Try 'ps --help <simple|list|output|threads|misc|all>'
or 'ps --help <s|l|o|t|m|a>'
for additional help text.

For more details see ps(1).
tsec@ubuntu:~/Desktop/S11-14-22-06$ ps -eo pid,ppid,cnd,%mem,%cpu --sort=-%mem | head
error: garbage option

Usage:
ps [options]

Try 'ps --help <simple|list|output|threads|misc|all>'
or 'ps --help <s|l|o|t|m|a>'
for additional help text.

For more details see ps(1).
tsec@ubuntu:~/Desktop/S11-14-22-06$ ps -eo pid,ppid,cnd,%mem,%cpu --sort=-%mem|head
error: garbage option

Usage:
ps [options]

Try 'ps --help <simple|list|output|threads|misc|all>'
or 'ps --help <s|l|o|t|m|a>'
for additional help text.

For more details see ps(1).
tsec@ubuntu:~/Desktop/S11-14-22-06$ ps -eo pid,ppid,cnd,%mem,%cpu --sort=-%mem|head
PID PPID CND %MEM %CPU
1855 1734 /usr/bin/gnome-shell 4.9 0.9
2213 1734 /usr/bin/gnome-software --g 3.5 0.0
862 825 /usr/bin/gnome-shell 3.5 0.1
1725 1723 /usr/lib/xorg/Xorg vt2 -dis 2.4 0.6
2082 1734 nautilus-desktop 2.3 0.0
2231 1 /usr/lib/fwupd/fwupd 2.0 0.0
2086 1700 /usr/lib/evolution/evolutio 1.6 0.0
1332 1 /usr/lib/packagekit/package 1.6 0.0
2117 2086 /usr/lib/evolution/evolutio 1.5 0.0
tsec@ubuntu:~/Desktop/S11-14-22-06$
```

Conclusion: Thus, we have successfully studied and implemented shell scripting languages

Name: Meet Brijwani

Roll No: 14

Batch: S11

Experiment No 3 : Linux- API

Aim: Implement Basic Commands of Linux like ls, cp, mv using Kernel APIs.

Theory:

1) cp → cp - r → copies the file cp - backups → It copies the backups of the destination file in the source folder.

cp - i → It asks for confirmation of the user whether yes or no and gives warning to the user before overriding the destination file.

cp -ie.txt b.txt

cp - f → Unable to open destination file for writing because the user has not allowed the writing operation .By writing the command, the destination file is deleted and the content is copied from source to destination.

cp - v → it is used for which command is running in backup.

cp - p → It preserves the attributes of files such as last date modification time, time of last access owners and the file permissions.

Syntax : cp-p source file destination file.

cp - l → To create a hand link file

2) mv → (move or rename). mv.i → Ask the user for confirmation before moving the file.

`mv - f` → It provides the protection and overrides destination file forcefully and deletes the source file.

`mv - f a.txt b.txt mv - n` → It takes/prevents the existing file from being overridden.

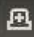




`mv - n a.txt b.txt mv - b` → To take the backup of an existing file

Output:

```
File Edit View Search Terminal Help
tsec10@ubuntu:~$ ps -aux
#167-18.04.1-Ubuntu SMP Wed May 24 00:51:42 UTC 2023
tsec10@ubuntu:~$ ps -aux
Linux ubuntu 5.4.0-150-generic #167-18.04.1-Ubuntu SMP Wed May 24 00:51:42 UTC 2023 x86_64 x86_64 GNU/Linux
tsec10 2214 1.0 3.6 1267372 144828 ? Sll 19:32 0:02 /usr/bin/gnome-software --application-service
root 2223 0.1 2.0 623192 80568 ? Ssl 19:32 0:00 /usr/lib/fwupd/fwupd
tsec10 2264 0.3 1.2 625444 51668 ? Sl 19:32 0:00 /usr/bin/nautilus --application-service
root 2282 0.0 0.0 0 0 ? I 19:32 0:00 [kworker/0:4-cgr]
root 2283 0.0 0.0 0 0 ? I 19:32 0:00 [kworker/0:5]
tsec10 2293 0.2 0.9 791688 36292 ? Ssl 19:33 0:00 /usr/lib/gnome-terminal/gnome-terminal-server
tsec10 2347 0.0 0.1 22452 4800 pts/0 Ss 19:33 0:00 bash
tsec10 2359 0.0 0.6 586872 25208 tty2 Sl+ 19:33 0:00 update-notifier
tsec10 2420 0.0 0.0 39672 3712 pts/0 R+ 19:35 0:00 ps -aux
tsec10 2421 0.0 0.0 7516 824 pts/0 S+ 19:35 0:00 tail
tsec10@ubuntu:~$ ps -aux|tail -n15
tsec10 2168 0.0 0.1 187908 5128 ? Sl 19:32 0:00 /usr/lib/dconf/dconf-service
tsec10 2168 0.1 1.5 1129008 62000 ? Sl 19:32 0:00 /usr/lib/evolution/evolution-calendar-factory-subprocess --factory all --bus-name org.gnome.evolution.dataserver.Subprocess.Backend.Calendar
arx2137x2 --own-path /org/gnome/evolution/dataserver/Subprocess/Backend/Calendar/2137/2
tsec10 2176 0.0 0.6 725716 24588 ? Ssl 19:32 0:00 /usr/lib/evolution/evolution-addressbook-factory
tsec10 2193 0.0 0.6 1075640 25204 ? Sl 19:32 0:00 /usr/lib/evolution/evolution-addressbook-factory-subprocess --factory all --bus-name org.gnome.evolution.dataserver.Subprocess.Backend.AddressBookX2178x2 --own-path /org/gnome/evolution/dataserver/Subprocess/Backend/AddressBook/2178/2
tsec10 2208 0.0 0.1 197368 5628 ? Ssl 19:32 0:00 /usr/lib/gvfs/gvfsd-metadata
tsec10 2214 0.9 3.6 1267372 144828 ? Sll 19:32 0:02 /usr/bin/gnome-software --application-service
root 2223 0.1 2.0 623192 80568 ? Ssl 19:32 0:00 /usr/lib/fwupd/fwupd
tsec10 2264 0.3 1.2 625444 51668 ? Sl 19:32 0:00 /usr/bin/nautilus --application-service
root 2282 0.0 0.0 0 0 ? I 19:32 0:00 [kworker/0:4-cgr]
root 2283 0.0 0.0 0 0 ? I 19:32 0:00 [kworker/0:5]
tsec10 2293 0.2 0.9 791688 36344 ? Ssl 19:33 0:00 /usr/lib/gnome-terminal/gnome-terminal-server
tsec10 2347 0.0 0.1 22452 4804 pts/0 Ss 19:33 0:00 bash
tsec10 2359 0.0 0.6 586872 25208 tty2 Sl+ 19:33 0:00 update-notifier
tsec10 2422 0.0 0.0 39672 3700 pts/0 R+ 19:36 0:00 ps -aux
tsec10 2423 0.0 0.0 7516 828 pts/0 S+ 19:36 0:00 tail -n15
tsec10@ubuntu:~$ ps -p55
PID TTY TIME CMD
2347 pts/0 00:00:00 bash
tsec10@ubuntu:~$ sudo apt install htop
[sudo] password for tsec10:
1
Sorry, try again.
[sudo] password for tsec10:
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
 fonts-liberation2 fonts-opensymbol gir1.2-gst-plugins-base-1.0 gir1.2-gudev-1.0 gir1.2-udisks-2.0
 grilo-plugins-0.3-base gstreamer1.0-gtk3 libboost-date-time1.65.1 libboost-filesystem1.65.1 libboost-iostreams1.65.1
 libboost-locale1.65.1 libcdt-0.1-1 libclucene-contribs1v5 libclucene-core1v5 libcmis-0.5-5v5 libcolamd2 libdazzle-1.0-0
 libe-book-0.1-1 libedataserverui-1.2-2 libeot0 libepubgen-0.1-1 libetonyek-0.1-1 libevent-2.1-6 libexiv2-14
 libfreerdp-client2-2 libfreerdp2-2 libgc1c2 libgee-0.8-2 libgexiv2-2 libgom-1.0-0 libgpgmepp6 libgpod-common libgpod4
 liblangtag-common liblangtag1 liblirc-client0 liblua5.3-0 libmediaart-2.0-0 libmsspub-0.1-1 libodfgen-0.1-1 libqgwing2v5
 libraw16 librevenge-0.0-0 libsguttl5-2 libssh-4 libsuitesparseconfig5 libvncclient1 libwinpr2-2 libxapian30 libxmlsec1-nss
 linux-hwe-5.4-headers-5.4.0-137 linux-hwe-5.4-headers-5.4.0-84 lp-solve media-player-info python3-mako python3-markupsafe
 syslinux syslinux-common syslinux-legacy usb-creator-common
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
 htop
0 upgraded, 1 newly installed, 0 to remove and 53 not upgraded.
Need to get 80.0 kB of archives.
After this operation, 221 kB of additional disk space will be used.
Get:1 http://us.archive.ubuntu.com/ubuntu bionic/main amd64 htop amd64 2.1.0-3 [80.0 kB]
Fetched 80.0 kB in 1s (84.8 kB/s)
Selecting previously unselected package htop.
(Reading database ... 186880 files and directories currently installed.)
Preparing to unpack .../htop_2.1.0-3_amd64.deb ...
Unpacking htop (2.1.0-3) ...
Setting up htop (2.1.0-3) ...
Processing triggers for desktop-file-utils (0.23-1ubuntu3.18.04.2) ...
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...
Processing triggers for gnome-menus (3.13.3-11ubuntu1.1) ...
Processing triggers for mime-support (3.60ubuntu1) ...
tsec10@ubuntu:~$ ps -eo pid,ppid,cmd,%mem,%cpu --sort=-%mem|head
PID PPID CMD %MEM %CPU
1901 1778 /usr/bin/gnome-shell 4.6 1.5
2214 1749 /usr/bin/gnome-software --g 3.6 0.3
1021 1015 /usr/bin/gnome-shell 3.5 0.2
2106 1778 nautilus-desktop 2.7 0.1
2223 1 /usr/lib/fwupd/fwupd 2.0 0.0
1769 1767 /usr/lib/xorg/Xorg vt2 -dis 1.9 0.6
2137 1749 /usr/lib/evolution/evolutio 1.6 0.0
1281 1 /usr/lib/packagekit/package 1.6 0.5
2168 2137 /usr/lib/evolution/evolutio 1.5 0.0
tsec10@ubuntu:~$
tsec10@ubuntu:~$
```

```
tsec10@ubuntu:~$ sudo apt install htop
[sudo] password for tsec10:
1
Sorry, try again.
[sudo] password for tsec10:
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
 fonts-liberation2 fonts-opensymbol gir1.2-gst-plugins-base-1.0 gir1.2-gstreamer-1.0 gir1.2-gudev-1.0 gir1.2-udisks-2.0
 grilo-plugins-0.3-base gstreamer1.0-gtk3 libboost-date-time1.65.1 libboost-filesystem1.65.1 libboost-iostreams1.65.1
 libboost-locale1.65.1 libcdt-0.1-1 libclucene-contribs1v5 libclucene-core1v5 libcmis-0.5-5v5 libcolamd2 libdazzle-1.0-0
 libe-book-0.1-1 libedataserverui-1.2-2 libeot0 libepubgen-0.1-1 libetonyek-0.1-1 libevent-2.1-6 libexiv2-14
 libfreerdp-client2-2 libfreerdp2-2 libgc1c2 libgee-0.8-2 libgexiv2-2 libgom-1.0-0 libgpgmepp6 libgpod-common libgpod4
 liblangtag-common liblangtag1 liblirc-client0 liblua5.3-0 libmediaart-2.0-0 libmsspub-0.1-1 libodfgen-0.1-1 libqgwing2v5
 libraw16 librevenge-0.0-0 libsguttl5-2 libssh-4 libsuitesparseconfig5 libvncclient1 libwinpr2-2 libxapian30 libxmlsec1-nss
 linux-hwe-5.4-headers-5.4.0-137 linux-hwe-5.4-headers-5.4.0-84 lp-solve media-player-info python3-mako python3-markupsafe
 syslinux syslinux-common syslinux-legacy usb-creator-common
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
 htop
0 upgraded, 1 newly installed, 0 to remove and 53 not upgraded.
Need to get 80.0 kB of archives.
After this operation, 221 kB of additional disk space will be used.
Get:1 http://us.archive.ubuntu.com/ubuntu bionic/main amd64 htop amd64 2.1.0-3 [80.0 kB]
Fetched 80.0 kB in 1s (84.8 kB/s)
Selecting previously unselected package htop.
(Reading database ... 186880 files and directories currently installed.)
Preparing to unpack .../htop_2.1.0-3_amd64.deb ...
Unpacking htop (2.1.0-3) ...
Setting up htop (2.1.0-3) ...
Processing triggers for desktop-file-utils (0.23-1ubuntu3.18.04.2) ...
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...
Processing triggers for gnome-menus (3.13.3-11ubuntu1.1) ...
Processing triggers for mime-support (3.60ubuntu1) ...
tsec10@ubuntu:~$ ps -eo pid,ppid,cmd,%mem,%cpu --sort=-%mem|head
PID PPID CMD %MEM %CPU
1901 1778 /usr/bin/gnome-shell 4.6 1.5
2214 1749 /usr/bin/gnome-software --g 3.6 0.3
1021 1015 /usr/bin/gnome-shell 3.5 0.2
2106 1778 nautilus-desktop 2.7 0.1
2223 1 /usr/lib/fwupd/fwupd 2.0 0.0
1769 1767 /usr/lib/xorg/Xorg vt2 -dis 1.9 0.6
2137 1749 /usr/lib/evolution/evolutio 1.6 0.0
1281 1 /usr/lib/packagekit/package 1.6 0.5
2168 2137 /usr/lib/evolution/evolutio 1.5 0.0
tsec10@ubuntu:~$
tsec10@ubuntu:~$
```



```
Open ▾  copy.c Save    

#include<stdlib.h>
int main()
{
FILE *fptr1, *fptr2;
char filename [100], c;
printf("Enter the filename to open for reading \n");
scanf("%s", filename); // Open one file for reading

fptr1 = fopen(filename, "r");

if(fptr1 ==NULL)
{printf("Cannot open file %s in", filename); exit(0);}
printf("Enter the filename to open for writing \n");
scanf("%s", filename);
// Open another file for writing

fptr2= fopen(filename, "w"); if (fptr2 == NULL)
{printf("Cannot open file %s\n", filename); exit(0);} // Read contents from file
c = fgetc(fptr1); while (c != EOF) {fputc(c, fptr2);
c = fgetc(fptr1);}printf("\nContents copied to %s", filename);
fclose(fptr1); fclose(fptr2); return 0;
}
```

```
tsec@ubuntu:~$ gcc -o copy.out copy.c
tsec@ubuntu:~$ ./copy.out
Enter the filename to open for reading
example.txt
Enter the filename to open for writing
example1.txt
Contents copied to example1.txt
```

```
tsec@ubuntu:~$ cat example1.txt
Hello!!!
```

Conclusion: Thus we have successfully studied and implemented various basic commands of Linux like ls, cp and mv using kernel APIs

Name: Meet Brijwani

Roll no: 14

Batch: S11

Experiment 4: Linux - Process

Aim:

- a. To create a child process in Linux using the fork system call. From the child process obtain the process ID of both child and parent by using getpid and getppid system call.
- b. Explore wait and waitpid before termination of process.

Theory:

1] System call

When a program is in user mode and requires access to Ram or hardware resources, it must ask the kernel to provide access to that resource. This is done via something called a system call. When a program makes a system call, the mode is switched from user mode to Kernel mode. This is called context switch. The kernel provides the resources which the program requested. System calls are made by user level programmes in following cases:

Creating, opening, closing and delete files in the system

Creating and managing new processes

Creating a connection in the network, sending and receiving packets

Requesting access to a hardware device like a mouse or a printer

2] Fork ()

The fork system call is used to create processes. When a process makes a fork().call, an exact copy of the process is created.

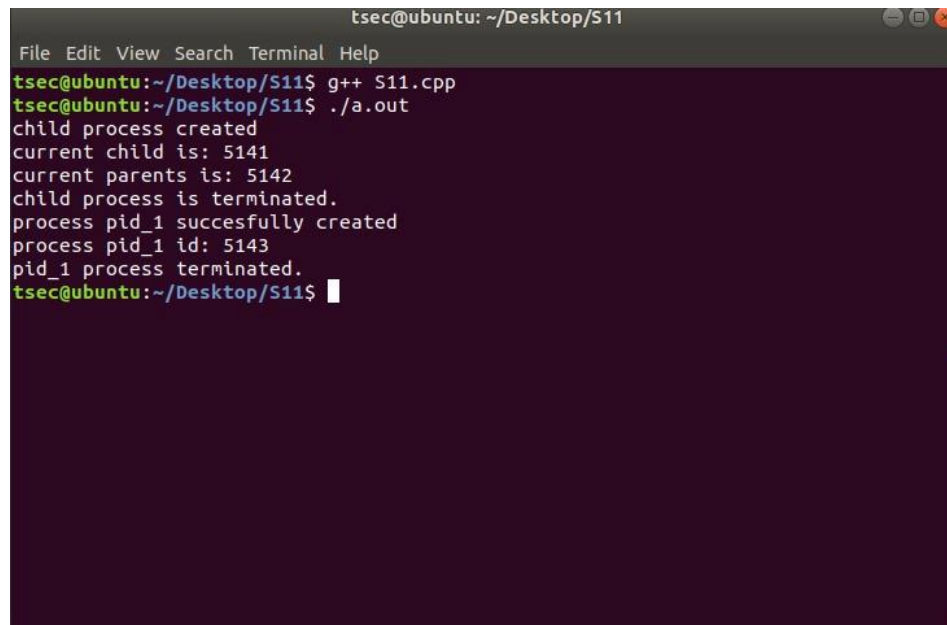
There are now two processes, one being the parent process and the other being the child process. The process which called fork().call is the parent process and the process which is created newly is called child process. The child process will be exactly the same as the parent. The process state of the parent i.e the address space, variables, open files etc is copied into the child process. The change of values in the parent process doesn't affect the child and vice versa.

Code:

```
#include <iostream>
#include<sys/wait.h>
#include<unistd.h>

using namespace std;
int wait_func()
{
    int pid_1 = fork();
    if(pid_1==0)
    {
        cout<<"process pid_1 succesfully
        created"<<"\n";    cout<<"process pid_1 id: "<<
        getpid()<<"\n";    exit(0); }
        waitpid(pid_1, NULL,0);
        cout << "pid_1 process terminated."<<"\n";
    return 0;
}
int main()
{
    int pid=fork();
    if(pid==0)
    {
        cout<<"child process created"<<"\n";
        cout<<"current child is: "<<getppid()<<"\n";
        cout<<"current parents is: "<<getpid()<<"\n";
        exit(0); } wait(NULL); cout<<"child process
        is terminated."<<"\n";
        wait_func();
        return 0;
    }
```

Output:



```
tsec@ubuntu: ~/Desktop/S11
File Edit View Search Terminal Help
tsec@ubuntu:~/Desktop/S11$ g++ S11.cpp
tsec@ubuntu:~/Desktop/S11$ ./a.out
child process created
current child is: 5141
current parents is: 5142
child process is terminated.
process pid_1 succesfully created
process pid_1 id: 5143
pid_1 process terminated.
tsec@ubuntu:~/Desktop/S11$
```

Conclusion:

Thus, we have successfully studied and implemented how to create child process in Linux using fork() system calls.

Name: Meet Brijwani

Roll no: 14

Batch: S-11

Aim : : a) To study and implement non preemptive scheduling algorithm FCFS.

b) To study and implement preemptive scheduling algorithm SRTF.

Theory :

Non preemptive algorithm :

FCFS, SJF 1] FCFS : First Come

First Serve

It is a non-preemptive scheduling algorithm and the criteria for this is the arrival time of the

process CPU is allotted to the process that requires it first. Jobs arriving later are placed at the end of the queue.

FCFS (Example)

Process	Duration	Oder	Arrival Time
P1	24	1	0
P2	3	2	0
P3	4	3	0

Gantt Chart :



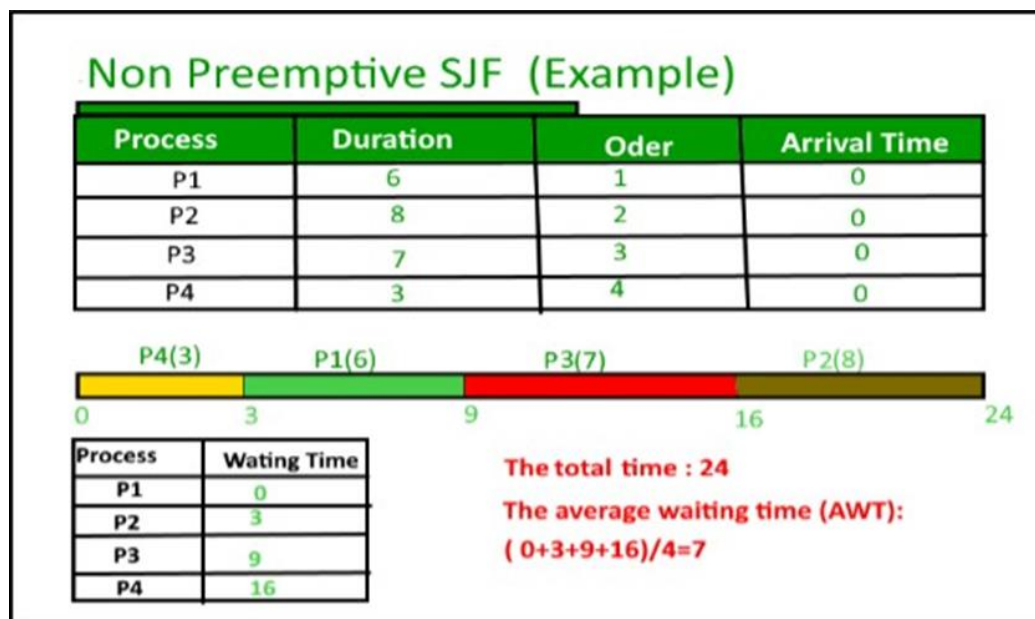
P1 waiting time : 0
P2 waiting time : 24
P3 waiting time : 27

The Average waiting time :
 $(0+24+27)/3 = 17$

2] SJF : shortest job first

This is a non preemptive scheduling algorithm, which associates with each process the length of the processes next CPU first. Criteria : Burst Time.

This algorithm assigns processes according to BT if BT of two processes is the same we use FCFS scheduling criteria.



3] Priority Scheduling

This is a non preemptive scheduling algorithm. Each process here has a priority that is either assigned already or externally done.

Process	Burst Time	Priority
P1	10	2
P2	5	0
P3	8	1

P1	P3	P2
-----------	-----------	-----------

0 10 18 23

Preemptive Scheduling Algorithm : SRTF/STRN

Shortest Remaining Time First / Shortest Remaining Time Next scheduling.

It is a preemptive SJF Algorithm. The choice arrives when a new process arrives as the ready

queue. While a previous process is still executing. The next CPU burst if the newly arrived

process may be shorter than what is left of the currently executing process. A preemptive SJF

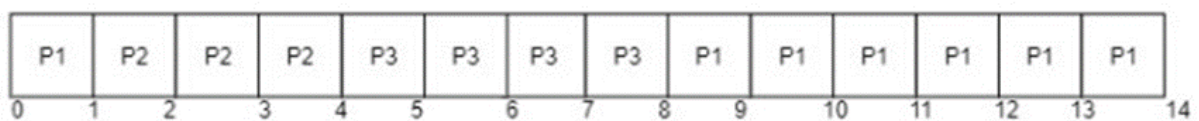
will preempt the current executing process and not allow the currently running process to finish

its CPU burst.

Round Robin is also one preemptive algorithm.

Process	Burst Time	Arrival Time
P1	7	0
P2	3	1
P3	4	3

The Gantt Chart for SRTF will be:



Implementation with code FCFS:

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int n;
    printf("Enter no. of processes : ");
    scanf("%d", &n);

    int *at = (int *)malloc(n * sizeof(int));
    int *bt = (int *)malloc(n * sizeof(int));
    int *ct = (int *)malloc(n * sizeof(int));
    int *wt = (int *)malloc(n * sizeof(int));
    int *tat = (int *)malloc(n * sizeof(int));

    printf("Enter arrival and burst times of the processes:\n");
    for (int i = 0; i < n; i++) {
        scanf("%d%d", &at[i], &bt[i]);
    }
}
```

```

int total_tat = 0;
int t = at[0];
for (int i = 0; i < n; i++) {
    t += bt[i];
    ct[i] = t;
    tat[i] = ct[i] - at[i];
    total_tat += tat[i];
    wt[i] = ct[i] - bt[i];
}

printf("  AT   BT   CT   TAT   WT\n");
for (int i = 0; i < n; i++) {
    printf("%6d%6d%6d%7d%6d\n", at[i], bt[i], ct[i], tat[i], wt[i]);
}

printf("Average TAT = %.2f\n", (float)total_tat / n);

free(at);
free(bt);
free(ct);
free(wt);
free(tat);

return 0;
}

```

Output :

```

Enter no. of processes : 3
Enter arrival and burst times of the processes:
0 5
1 3
2 8

```

	AT	BT	CT	TAT	WT
0	0	5	5	5	0
1	1	3	8	7	5
2	2	8	16	14	8

```

Average TAT = 8.67

```

Implementation with code SRTF:

```
#include <stdio.h>
#include <stdlib.h>
#include <limits.h>

struct Process {
    int pid; // Process ID
    int bt; // Burst Time
    int art; // Arrival Time
};

void findWaitingTime(struct Process proc[], int n, int wt[]) {
    int rt[n];
    for (int i = 0; i < n; i++)
        rt[i] = proc[i].bt;

    int complete = 0, t = 0, minm = INT_MAX;
    int shortest = 0, finish_time;
    int check = 0;
```

```

while (complete != n) {
    for (int j = 0; j < n; j++) {
        if ((proc[j].art <= t) && (rt[j] < minm) && rt[j] > 0) {
            minm = rt[j];
            shortest = j;
            check = 1;
        }
    }

    if (check == 0) {
        t++;
        continue;
    }

    rt[shortest]--;
    minm = rt[shortest];
    if (minm == 0)
        minm = INT_MAX;

    if (rt[shortest] == 0) {
        complete++;
        check = 0;
        finish_time = t + 1;
        wt[shortest] = finish_time - proc[shortest].bt - proc[shortest].art;
        if (wt[shortest] < 0)
            wt[shortest] = 0;
    }
    t++;
}
}

```

```

void findTurnAroundTime(struct Process proc[], int n, int wt[], int tat[]) {
    for (int i = 0; i < n; i++)
        tat[i] = proc[i].bt + wt[i];
}

```

```

void findavgTime(struct Process proc[], int n) {

```

```

int *wt = (int *)malloc(n * sizeof(int));
int *tat = (int *)malloc(n * sizeof(int));
int total_wt = 0, total_tat = 0;

findWaitingTime(proc, n, wt);
findTurnAroundTime(proc, n, wt, tat);

printf(" P\t\tBT\t\tWT\t\tTAT\t\t\n");
for (int i = 0; i < n; i++) {
    total_wt += wt[i];
    total_tat += tat[i];
    printf(" %d\t\t%d\t\t%d\t\t%d\n", proc[i].pid, proc[i].bt, wt[i], tat[i]);
}

printf("\nAverage waiting time = %.2f\n", (float)total_wt / (float)n);
printf("Average turn around time = %.2f\n", (float)total_tat / (float)n);

free(wt);
free(tat);
}

int main() {
    int n;
    printf("Enter the number of processes: ");
    scanf("%d", &n);

    struct Process *proc = (struct Process *)malloc(n * sizeof(struct
Process));
    printf("Enter arrival time and burst time for each process:\n");
    for (int i = 0; i < n; i++) {
        printf("Process %d: ", i + 1);
        scanf("%d %d", &proc[i].art, &proc[i].bt);
        proc[i].pid = i + 1;
    }

    findavgTime(proc, n);

```



```
    free(proc);  
  
    return 0;  
}
```

Output:

```
Enter the number of processes: 5  
Enter arrival time and burst time for each process:  
Process 1: 2 6  
Process 2: 5 2  
Process 3: 1 8  
Process 4: 0 3  
Process 5: 4 4  
P      BT      WT      TAT  
1      6      7      13  
2      2      0      2  
3      8      14     22  
4      3      0      3  
5      4      2      6  
  
Average waiting time = 4.60  
Average turn around time = 9.20
```

Conclusion : Thus we have successfully implemented preemptive and non preemptive scheduling algorithms.

Name: Meet Brijwani

Roll no: 14

Batch: S11

EXP 6:

Aim : Write a C program to implement solution of Producer consumer problem through Semaphore.

Theory :

In concurrent programming, concurrency represents a pivotal concept necessary to comprehend fully how such systems operate. Among the various challenges encountered by practitioners working with these systems stands out the producer-consumer problem - one of the most renowned synchronization issues. In this text, our objective consists of analyzing this topic and highlighting its significance for concurrent computing while also examining possible solutions rooted within C.

Introduction

In concurrent systems, multiple threads or processes may access shared resources simultaneously. The producer-consumer problem involves two entities: producers that generate data or tasks, and consumers that process or consume the generated data. The challenge lies in ensuring that producers and consumers synchronize their activities to avoid issues like race conditions or resource conflicts.

Understanding the Producer-Consumer Problem

Problem Statement

One possible definition of the producer-consumer problem involves two main groups: producers of data who store their work in a communal space called the buffer, and processors (consumers) of that content saved in said space. These

people use their expertise around gathered items within this temporary holding scenario to analyze it comprehensively before delivering insightful results.

Synchronization Requirements

Achieving resolution of the producer-consumer conundrum will necessarily involve implementing techniques for synchronized collaboration among all stakeholders involved. The integration of optimal synchronization protocols is fundamental in avoiding scenarios where device buffers are either overloaded by producing units or depleted by consuming ones.

Implementing the Producer-Consumer Problem in C

Shared Buffer

In C, a shared buffer can be implemented using an array or a queue data structure. The buffer should have a fixed size and support operations like adding data (producer) and retrieving data (consumer).

Synchronization Techniques

Several synchronization techniques can be used to solve the producer-consumer problem in C, including –

- Mutex and Condition Variables – Mutexes provide mutual exclusion to protect critical sections of code, while condition variables allow threads to wait for specific conditions to be met before proceeding.

- Semaphores – Semaphores can be used to control access to the shared buffer by tracking the number of empty and full slots.

- Monitors – Monitors provide a higher-level abstraction for synchronization and encapsulate shared data and the operations that can be performed on it.

Solutions to the Producer-Consumer Problem in C

Bounded Buffer Solution

One common solution to the producer-consumer problem is the bounded buffer solution. It involves using a fixed-size buffer with synchronization mechanisms to ensure that producers and consumers cooperate correctly. The capacity for

item production is limited by the buffer size making it crucial to factor in this specification so as not to exceed the available space in the buffer.

Producer and Consumer Threads

In C, the producer and consumer activities can be implemented as separate threads. Each producer thread generates data and adds it to the shared buffer, while each consumer thread retrieves data from the buffer and processes it. Synchronization mechanisms are used to coordinate the activities of the threads.

Handling Edge Cases

In real-world scenarios, additional considerations may be necessary. For example, if the producers generate data at a faster rate than the consumers can process, buffering mechanisms like blocking or dropping data may be required to prevent data loss or deadlock situations.

Code :

```
#include <stdio.h>
#include <stdlib.h>

// Initialize a mutex to 1 and Number of full slots as 0 and Number of
empty slots as size of buffer

int mutex = 1;
int full = 0;
int empty = 10, x = 0;

// Function to produce an item and add it to the buffer
void producer()
{
    // Decrease mutex value by 1 and Increase the number of full slots
```

```
by 1 and Decrease the number of empty slots by 1
```

```
--mutex;
```

```
++full;
```

```
--empty;
```

```
// Item produced
```

```
x++;
```

```
printf("\nProducer produces item %d",x);
```

```
// Increase mutex value by 1
```

```
++mutex;
```

```
}
```

```
// Function to consume an item and remove it from buffer
```

```
void consumer()
```

```
{
```

```
    // Decrease mutex value by 1 and Decrease the number of full slots  
by 1
```

```
--mutex;
```

```
--full;
```

```
// Increase the number of empty slots by 1
```

```
++empty;
```

```
printf("\nConsumer consumes item %d",x);
```

```
x--;
```

```
// Increase mutex value by 1
```

```
++mutex;
```

```
}
```

```
int main()
```

```

{

    int n, i; printf("\n1. Press 1 for Producer \n2. Press 2 for
    Consumer \n3.
Press 3 for Exit");

    for (i = 1; i > 0; i++) {

```

```

        printf("\nEnter your choice:");

        scanf("%d", &n);

        switch (n) {

            case 1:

                // If mutex is 1 and empty is non-zero, then it is possible to produce

                if ((mutex == 1)

                    && (empty != 0)) {

                    producer();

                }

                else { printf("Buffer is
                    full!");

                }

                break;

            case 2:

                // If mutex is 1 and full is non-zero, then it is possible to consume

                if ((mutex == 1) && (full != 0)) {

                    consumer();

                }

```

```
        else {  
            printf("Buffer is empty!");  
        }  
        break;  
    case 3:  
        exit(0);  
        break;  
    }  
}
```

Output :

Output

Clear

/tmp/QgcMuqYbY.o

1. Press 1 for Producer

2. Press 2 for Consumer

3. Press 3 for Exit

Enter your choice:2

Buffer is empty!

Enter your choice:1

Producer produces item 1

Enter your choice:1

Producer produces item 2

Enter your choice:1

Producer produces item 3

Enter your choice:1

Producer produces item 4

Enter your choice:1

Producer produces item 5

Enter your choice:1

Producer produces item 6

Enter your choice:1

```
Producer produces item 7
Enter your choice:1

Producer produces item 8
Enter your choice:1

Producer produces item 9
Enter your choice:1

Producer produces item 10
Enter your choice:1
Buffer is full!
Enter your choice:2

Consumer consumes item 10
Enter your choice:2

Consumer consumes item 9
Enter your choice:2

Consumer consumes item 8
Enter your choice:2

Consumer consumes item 7
Enter your choice:2
```

```
Consumer consumes item 6
Enter your choice:2

Consumer consumes item 5
Enter your choice:2

Consumer consumes item 4
Enter your choice:2

Consumer consumes item 3
Enter your choice:2

Consumer consumes item 2
Enter your choice:2

Consumer consumes item 1
Enter your choice:2
Buffer is empty!
Enter your choice:3
```

Conclusion : Thus we have successfully implemented Producer Consumer problem using Semaphore.

Name: Meet Brijwani

Roll no: 14

Batch: S11

EXP 7:

Aim : Process Management: Deadlock

- a. Write a program to demonstrate the concept of deadlock avoidance through Banker's Algorithm
- b. Write a program demonstrate the concept of Dining Philosopher's Problem

Theory : The banker's algorithm is a resource allocation and deadlock avoidance algorithm that tests for safety by simulating the allocation for the predetermined maximum possible amounts of all resources, then makes an "s-state" check to test for possible activities, before deciding whether allocation should be allowed to continue

Available

It is a 1-d array of size 'm' indicating the number of available resources of each type.

$\text{Available}[j] = k$ means there are 'k' instances of resource type R_j

Max

It is a 2-d array of size 'n*m' that defines the maximum demand of each process in a system.

$\text{Max}[i, j] = k$ means process P_i may request at most 'k' instances of resource type R_j .

Allocation

It is a 2-d array of size 'n*m' that defines the number of resources of each type currently allocated to each process.

$\text{Allocation}[i, j] = k$ means process P_i is currently allocated 'k' instances of resource type R_j

Need

It is a 2-d array of size 'n*m' that indicates the remaining resource need of each process.

Need [i, j] = k means process P_i currently needs 'k' instances of resource type R_j
Need [i, j] = Max [i, j] – Allocation [i, j]

Allocation specifies the resources currently allocated to process P_i and Need $_i$ specifies the additional resources that process P_i may still request to complete its task. Banker's algorithm consists of a Safety algorithm and a Resource request algorithm.

Banker's Algorithm

1. Active:= Running U Blocked;

for $k=1 \dots r$

New_request[k]:= Requested_resources[requesting_process, k];

2. Simulated_allocation:= Allocated_resources; for $k=1 \dots r$

//Compute projected allocation state

Simulated_allocation [requesting_process, k]:= Simulated_allocation [requesting_process, k] + New_request[k]; 3. feasible:= true; for

$k=1 \dots r$ // Check whether projected allocation state is feasible if Total_resources[k] < Simulated_total_alloc [k] then feasible:= false;

4. if feasible= true

then // Check whether projected allocation state is a safe allocation state

while set Active contains a process P_1 such that

For all k, Total_resources[k] – Simulated_total_alloc[k] >= Max_need [1

,k]-Simulated_
allocation[l, k]

Delete P1 from Active;

for k=1.....r

Simulated_ total_ alloc[k]:= Simulated_ total_ alloc[k]- Simulated_ allocation[l, k];

5. If set Active is empty then // Projected allocation state is a safe

allocation state for k=1....r // Delete the request from pending

requests Requested_ resources[requesting_ process, k]:=0; for

k=1....r // Grant the request

Allocated_ resources[requesting_ process, k]:= Allocated_
resources[requesting_ process, k] + New_ request[k];

Total_ alloc[k]:= Total_ alloc[k] + New_ request[k];

Safety Algorithm

The algorithm for finding out whether or not a system is in a safe state can be described as follows:

1) Let Work and Finish be vectors of length 'm' and 'n' respectively.

Initialize: Work = Available

Finish[i] = false; for i=1, 2, 3, 4....n

2) Find an i such that both a)

Finish[i] = false

b) Needi <= Work if no such i

exists goto step (4) 3) Work =

Work + Allocation[i] Finish[i]

= true goto step (2)

4) if Finish [i] = true for all i

then the system is in a safe state Safety Algorithm

Resource-Request Algorithm

Let $Request_i$ be the request array for process P_i . $Request_i[j] = k$ means process P_i wants k

instances of resource type R_j . When a request for resources is made by process P_i , the following actions are taken:

1) If $Request_i \leq Need_i$

Goto step (2) ; otherwise, raise an error condition, since the process has exceeded its maximum claim.

2) If $Request_i \leq Available$

Goto step (3); otherwise, P_i must wait, since the resources are not available. 3) Have the system pretend to have allocated the requested resources to process P_i by modifying the state as follows:

$Available = Available - Request_i$

$Allocation_i = Allocation_i + Request_i$

$Need_i = Need_i - Request_i$

Code : `// Banker's Algorithm`

```
#include <iostream>
```

```
using namespace std;
```

```
int main()
```

```
{
```

```
int n, m, i, j, k;
```

```
n = 5;
```

```
m = 3;
```

```
int alloc[5][3] = { { 0, 1, 0 },
```

```
{ 2, 0, 0 },
```

```
{ 3, 0, 2 },
```

```
{ 2, 1, 1 },
```

```
{ 0, 0, 2 } };
```

```
int max[5][3] = { { 7, 5, 3 },
```

```
{ 3, 2, 2 },
```

```
{ 9, 0, 2 },
```

```
{ 2, 2, 2 },
```

```
{ 4, 3, 3 } };
```

```
int avail[3] = { 3, 3, 2 };
```

```
int f[n], ans[n], ind = 0;
```

```
for (k = 0; k < n; k++) {  
f[k] = 0;
```



```
}

int need[n][m];

for (i = 0; i < n; i++) {
```

```
    for (j = 0; j < m; j++)
        need[i][j] = max[i][j] - alloc[i][j];
}

int y = 0;

for (k = 0; k < 5; k++) {
    for (i = 0; i < n; i++) {
        if (f[i] == 0) {

            int flag = 0;

            for (j = 0; j < m; j++) {
                if (need[i][j] > avail[j]){
                    flag = 1;
                    break;
                }
            }

            if (flag == 0) {
                ans[ind++] = i;

                for (y = 0; y < m; y++)
                    avail[y] += alloc[i][y];

                f[i] = 1;
            }
        }
    }
}
```

```

}

}

}

}

int flag = 1;

for(int i = 0;i<n;i++)

{

    if(f[i]==0)
    {

        flag = 0;

        cout << "The given sequence is not safe";

        break;

    }

}

if(flag==1)
{

    cout << "Following is the SAFE Sequence" << endl;

    for (i = 0; i < n - 1; i++)

        cout << " P" << ans[i] << " ->";

    cout << " P" << ans[n - 1] <<endl;

} return

(0);

}

```

Output :

```
#include <stdio.h>
```

Following is the SAFE Sequence

P1 -> P3 -> P4 -> P0 -> P2

=== Code Execution Successful ===

Code :

```
#include <stdlib.h>
```

```
#include <pthread.h>
```

```
#include <semaphore.h>
```

```
#define NUM_PHILOSOPHERS 5
```

```
#define NUM_CHOPSTICKS 5
```

```
void dine(int n);
```

```
pthread_t philosopher[NUM_PHILOSOPHERS];
```

```
pthread_mutex_t chopstick[NUM_CHOPSTICKS];
```

```
int main()
```

```
{

int i, status_message;

void *msg;

for (i = 1; i <= NUM_CHOPSTICKS; i++)

{

status_message = pthread_mutex_init(&chopstick[i], NULL);

if (status_message == -1)

{

printf("\n Mutex initialization failed");
exit(1);

}

}

for (i = 1; i <= NUM_PHILOSOPHERS; i++)
```

```
{

status_message = pthread_create(&philosopher[i], NULL, (void *)dine,
(int *)i);
```

```
if (status_message != 0)
{
printf("\n Thread creation error \n");
exit(1);
}
}

for (i = 1; i <= NUM_PHILOSOPHERS; i++)
{
status_message = pthread_join(philosopher[i], &msg);

if (status_message != 0)
{
printf("\n Thread join failed \n");
exit(1);
}
}

for (i = 1; i <= NUM_CHOPSTICKS; i++)
{
status_message = pthread_mutex_destroy(&chopstick[i]);

if (status_message != 0)
{
printf("\n Mutex Destroyed \n");
exit(1);
}
```

```
}

return 0;

}

void dine(int n)
{
    printf("\nPhilosopher %d is thinking ", n);

    pthread_mutex_lock(&chopstick[n]);

    pthread_mutex_lock(&chopstick[(n + 1) % NUM_CHOPSTICKS]);

    printf("\nPhilosopher %d is eating ", n);

    sleep(3);

    pthread_mutex_unlock(&chopstick[n]);

    pthread_mutex_unlock(&chopstick[(n + 1) % NUM_CHOPSTICKS]);
}
```

```
printf("\nPhilosopher %d Finished eating ", n);  
  
}
```

Output :

```
Philosopher 2 is thinking  
Philosopher 2 is eating  
Philosopher 3 is thinking  
Philosopher 5 is thinking  
Philosopher 5 is eating  
Philosopher 1 is thinking  
Philosopher 4 is thinking  
Philosopher 4 is eating  
Philosopher 2 Finished eating  
Philosopher 5 Finished eating  
Philosopher 1 is eating  
Philosopher 4 Finished eating  
Philosopher 3 is eating  
Philosopher 1 Finished eating  
Philosopher 3 Finished eating
```

Conclusion : Thus we have successfully implemented Banker's Algorithm and the concept of Dining Philosopher's Problem using C++.

Name: Meet Brijwani

Roll no: 14

Batch: S11

EXP 8:

Aim: Write a program to demonstrate the concept of dynamic partitioning placement algorithms i.e. Best Fit, First Fit, Worst Fit.

Theory:

In the operating system, the following are four common memory management techniques.

Single contiguous allocation: Simplest allocation method used by MS-DOS. All memory (except some reserved for OS) is available to a process.

Partitioned allocation: Memory is divided into different blocks or partitions. Each process is allocated according to the requirement.

Paged memory management: Memory is divided into fixed-sized units called page frames, used in a virtual memory environment.

Segmented memory management: Memory is divided into different segments (a segment is a logical grouping of the process' data or code). In this management, allocated memory doesn't have to be contiguous.

Most of the operating systems (for example Windows and Linux) use Segmentation with Paging. A process is divided into segments and individual segments have pages.

In Partition Allocation, when there is more than one partition freely available to accommodate a process's request, a partition must be selected. To choose a particular partition, a partition allocation method is needed. A partition allocation method is considered better if it avoids internal fragmentation.

When it is time to load a process into the main memory and if there is more than one free block of memory of sufficient size then the OS decides which free block to allocate.

There are different Placement Algorithm:

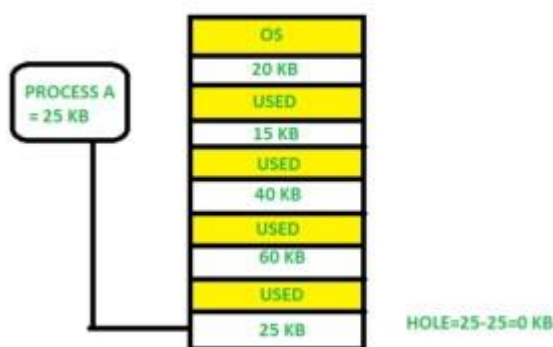
A. First Fit

- B. Best Fit
- C. Worst Fit
- D. Next Fit

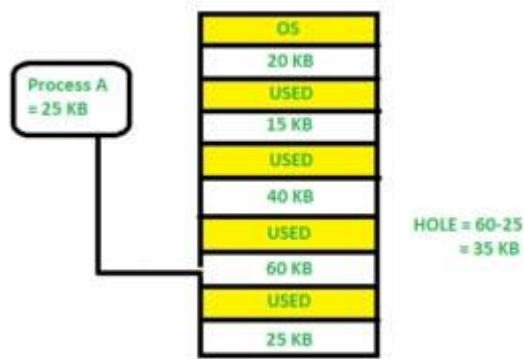
1. First Fit: In the first fit, the partition is allocated which is the first sufficient block from the top of Main Memory. It scans memory from the beginning and chooses the first available block that is large enough. Thus it allocates the first hole that is large enough.



2. Best Fit Allocate the process to the partition which is the first smallest sufficient partition among the free available partition. It searches the entire list of holes to find the smallest hole whose size is greater than or equal to the size of the process.



3. Worst Fit Allocate the process to the partition which is the largest sufficient among the freely available partitions available in the main memory. It is opposite to the best-fit algorithm. It searches the entire list of holes to find the largest hole and allocate it to process.



4. Next Fit: Next fit is similar to the first fit but it will search for the first sufficient partition from the last allocation point.

Is Best-Fit really best?

Although best fit minimises the wastage space, it consumes a lot of processor time for searching the block which is close to the required size. Also, Best-fit may perform poorer than other algorithms in some cases. For example, see the exercise below.

Sl.No.	Partition Allocation Method	Advantages	Disadvantages
1.	Fixed Partition	Simple, easy to use, no complex algorithms needed	Memory waste, inefficient use of memory resources
2.	Dynamic Partition	Flexible, more efficient, partitions allocated as required	Requires complex algorithms for memory allocation
3.	Best-fit Allocation	Minimizes memory waste, allocates smallest suitable partition	More computational overhead to find smallest split
4.	Worst-fit Allocation	Ensures larger processes have sufficient memory	May result in substantial memory waste

5.

First-fit
Allocation

Quick, efficient, less
computational work

Risk of memory
fragmentation

Comparison of Partition Allocation Methods: Exercise: Consider the requests from processes in given order 300K, 25K, 125K, and 50K. Let there be two blocks of memory available of size 150K followed by a block size 350K.

Which of the following partition allocation schemes can satisfy the above requests?

- A) Best fit but not first fit.
- B) First fit but not best fit.
- C) Both First fit & Best fit.
- D) neither first fit nor best fit.

Solution: Let us try all options.

Best Fit:

300K is allocated from a block of size 350K. 50 is left in the block.

25K is allocated from the remaining 50K block. 25K is left in the block.

125K is allocated from 150 K block. 25K is left in this block also.

50K can't be allocated even if there is 25K + 25K space available.

First Fit:

300K request is allocated from 350K block, 50K is left out.

25K is allocated from the 150K block, 125K is left out.

Then 125K and 50K are allocated to the remaining left out partitions.

So, the first fit can handle requests.

So option B is the correct choice.

Program:

a)First Fit:

```
#include<stdio.h>
```

```
// Function to allocate memory to
```

```
// blocks as per First fit algorithm
```

```
void firstFit(int blockSize[], int m, int processSize[], int n)
```

```
{ int i, j;
```

```

// Stores block id of the //
block allocated to a process
int allocation[n];

// Initially no block is assigned to any process
for(i = 0; i < n; i++)
{
    allocation[i] = -1;
}

// pick each process and find suitable blocks
// according to its size and assign to it
for (i = 0; i < n; i++) //here, n -> number of processes
{
    for (j = 0; j < m; j++) //here, m -> number of blocks
    {
        if (blockSize[j] >= processSize[i])
        {
            // allocating block j to the ith process
            allocation[i] = j;

            // Reduce available memory in this block.
            blockSize[j] -= processSize[i];

            break; //go to the next process in the queue
        }
    }
}

printf("\nProcess No.\tProcess Size\tBlock no.\n");
for (int i = 0; i < n; i++)
{ printf(" %i\t\t\t", i+1);
    printf("%i\t\t\t\t",
        processSize[i]);
    if (allocation[i] != -1)
        printf("%i", allocation[i] + 1);
}

```

```

        else
            printf("Not Allocated");
            printf("\n");
    }
}

// Driver code
int main()
{
    int m; //number of blocks in the memory
    int n; //number of processes in the input queue
    int blockSize[] = {100, 500, 200, 300, 600};
    int processSize[] = {212, 417, 112, 426};
    m = sizeof(blockSize) / sizeof(blockSize[0]);
    n = sizeof(processSize) / sizeof(processSize[0]);
    firstFit(blockSize, m, processSize, n);

    return 0 ;
}

```

b)Best Fit:

```
#include<stdio.h>
```

```

void main()
{
    int fragment[20],b[20],p[20],i,j,nb,np,temp,lowest=9999;
    static int barray[20],parray[20];
    printf("\n\t\t\tMemory Management Scheme - Best Fit");
    printf("\nEnter the number of blocks:");
    scanf("%d",&nb);
    printf("Enter the number of processes:");
    scanf("%d",&np);
    printf("\nEnter the size of the blocks:-\n");
    for(i=1;i<=nb;i++)
    {
        printf("Block no. %d:",i);
        scanf("%d",&b[i]);
    }
}

```

```

    }
    printf("\nEnter the size of the processes :-\n");
    for(i=1;i<=np;i++)
    {
        printf("Process no.:%d:",i);
        scanf("%d",&p[i]);
    }
    for(i=1;i<=np;i++)
    {
        for(j=1;j<=nb;j++)
        {
            if(barray[j]!=1)
            { temp=b[j]-p[i];
              if(temp>=0)
              if(lowest>temp)
              { parray[i]=j;
                lowest=temp;
              }
            }
        }
        fragment[i]=lowest;
        barray[parray[i]]=1;
        lowest=10000;
    }
    printf("\nProcess_no\tProcess_size\tBlock_no\tBlock_size\tFragment");
    for(i=1;i<=np && parray[i]!=0;i++)
    printf("\n%d\t%d\t%d\t%d\t%d",i,p[i],parray[i],b[parray[i]],fragment[i])
    ;
}

```

c) Worst Fit:

```
#include <stdio.h>
```

```

void implimentWorstFit(int blockSize[], int blocks, int processSize[], int
processes)
{
    // This will store the block id of the allocated block to a process
    int allocation[processes]; int occupied[blocks];

```

```

// initially assigning -1 to all allocation
indexes // means nothing is allocated
currently for(int i = 0; i < processes; i++){
allocation[i] = -1;
}

for(int i = 0; i < blocks; i++){
    occupied[i] = 0;
}

// pick each process and find suitable
blocks // according to its size and assign to it
for (int i=0; i < processes; i++)
{
int indexPlaced = -1;
for(int j = 0; j < blocks;
j++)
{
    // if not occupied and block size is large enough
    if(blockSize[j] >= processSize[i] && !occupied[j])
    {
        // place it at the first block fit to accommodate process
        if (indexPlaced == -1)
            indexPlaced = j;

        // if any future block is larger than the current block where //
        process is placed, change the block and thus indexPlaced
        else if (blockSize[indexPlaced] < blockSize[j])
            indexPlaced = j;
    }
}
}

// If we were successfully able to find block for the process
if (indexPlaced != -1)
{

```

```

        // allocate this block j to process p[i]
        allocation[i] = indexPlaced;

        // make the status of the block as occupied
        occupied[indexPlaced] = 1;

        // Reduce available memory for the block
        blockSize[indexPlaced] -= processSize[i];
    }
}

printf("\nProcess No.\tProcess Size\tBlock no.\n");
for (int i = 0; i < processes; i++)
{
    printf("%d \t\t\t %d \t\t\t", i+1, processSize[i]);
    if (allocation[i] != -1)
        printf("%d\n", allocation[i] + 1);
    else
        printf("Not Allocated\n");
}
}

// Driver code
int main()
{
    int blockSize[] = {100, 50, 30, 120, 35}; int processSize[] =
    {40, 10, 30, 60}; int blocks =
    sizeof(blockSize)/sizeof(blockSize[0]); int processes =
    sizeof(processSize)/sizeof(processSize[0]);
    implimentWorstFit(blockSize, blocks, processSize,
    processes);

    return 0;
}

```


Output:

a)First Fit:

```
PS C:\Meet\COLLEGE\LAB ASS\Aoa> cd "c:\Meet\COLLEGE\LAB ASS\Aoa\" ; if ($?) { gcc firstfit.c -o firstfit } ; if ($?) { .\firstfit }

Process No.    Process Size    Block no.
1              212            2
2              417            5
3              112            2
4              426            Not Allocated
PS C:\Meet\COLLEGE\LAB ASS\Aoa>
```

b)Best Fit:

```
PS C:\Meet\COLLEGE\LAB ASS\Aoa> cd "c:\Meet\COLLEGE\LAB ASS\Aoa\" ; if ($?) { gcc bestfit.c -o bestfit } ; if ($?) { .\bestfit }

Memory Management Scheme - Best Fit
Enter the number of blocks:5
Enter the number of processes:3

Enter the size of the blocks:-
Block no.1:10
Block no.2:20
Block no.3:30
Block no.4:30
Block no.5:40

Enter the size of the processes :-
Process no.1:7
Process no.2:9
Process no.3:2

Process_no    Process_size    Block_no    Block_size    Fragment
1             7              1           10           3
2             9              2           20           11
3             2              3           30           28
PS C:\Meet\COLLEGE\LAB ASS\Aoa>
```

c)Worst Fit:

```
PS C:\Meet\COLLEGE\LAB ASS\Aoa> cd "c:\Meet\COLLEGE\LAB ASS\Aoa\" ; if ($?) { gcc worstfit.c -o worstfit } ; if ($?) { .\worstfit }

Process No.    Process Size    Block no.
1              40            4
2              10            1
3              30            2
4              60            Not Allocated
PS C:\Meet\COLLEGE\LAB ASS\Aoa>
```

Conclusion: Thus we have successfully implemented dynamic partitioning placement algorithms.

Experiment no: 09.

Name: Meet Brijwani

Roll no: 14

Batch: S11

Aim: Write a program in C demonstrate the concept of page replacement policies for handling page faults eg: FIFO, LRU

Theory:

Page replacement policies are an integral part of virtual memory management in operating systems. When a program accesses data or instructions, the operating system loads the corresponding pages into physical memory (RAM) from secondary storage (usually a hard disk). However, physical memory has limited capacity, so not all pages can reside in memory simultaneously. When a program requests a page that is not in memory, a page fault occurs, triggering the need for a page replacement policy to determine which page to evict from memory to make room for the new one. Two commonly used page replacement policies are First-In-First-Out (FIFO) and Least Recently Used (LRU). Let's delve into these policies and understand their mechanisms and trade-offs.

First-In-First-Out (FIFO):

FIFO is one of the simplest page replacement algorithms. It evicts the oldest page in memory, based on the assumption that the page that has been in memory the longest is least likely to be needed in the near future.

Mechanism:

When a page fault occurs and memory is full, the operating system selects the page that entered memory earliest (the oldest page) for replacement.

The selected page is evicted from memory, and the new page is loaded in its place.

The page table is updated accordingly.

Advantages:

Simplicity: FIFO is easy to implement and understand.

Low Overhead: The overhead of maintaining data structures is minimal.

Disadvantages:

Belady's Anomaly: FIFO can suffer from Belady's anomaly, where increasing the number of frames can actually increase the number of page faults.

Poor Performance: FIFO does not consider the access history of pages, leading to suboptimal performance in many cases, especially when the access patterns are irregular.

Least Recently Used (LRU):

LRU is based on the principle that the page that has not been accessed for the longest time is least likely to be used in the near future.

Mechanism:

LRU keeps track of the time of the last access for each page.

When a page fault occurs, the operating system selects the page that was least recently accessed for replacement.

The selected page is evicted from memory, and the new page is loaded in its place.

The page table is updated accordingly, and the access time of the new page is recorded.

Advantages:

Optimality: LRU provides better performance than FIFO in terms of reducing the number of page faults in many scenarios.

Flexibility: LRU can adapt to varying access patterns by considering the access history of pages.

Disadvantages:

Implementation Complexity: Implementing an efficient LRU algorithm requires maintaining a data structure to track the access times of pages, which can be resource-intensive.

High Overhead: The overhead of maintaining access times for each page can be significant, especially in systems with a large number of pages.

Comparison and Trade-offs:

Optimality: LRU is generally considered more optimal than FIFO because it takes into account the actual access history of pages rather than just the order of arrival. However, implementing a true LRU algorithm can be complex and resource-intensive.

Overhead: FIFO has lower overhead compared to LRU since it does not require tracking access times for each page. However, this simplicity comes at the cost of potentially poorer performance.

Belady's Anomaly: FIFO can suffer from Belady's anomaly, where increasing the number of frames can paradoxically increase the number of page faults. LRU does not suffer from this anomaly.

Adaptability: LRU is more adaptable to varying access patterns since it considers the actual access history of pages. FIFO, on the other hand, may perform poorly in scenarios with irregular access patterns.

Program:

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#define MAX_FRAMES 3 // Maximum number of page frames
```

```
#define MAX_PAGES 10 // Maximum number of pages in reference  
string
```

```
// Function prototypes void fifo(int pages[], int n, int frames); void lru(int  
pages[], int n, int frames);
```

```
int main() {    int pages[MAX_PAGES];  
    int n, frames;
```

```
    printf("Enter the number of pages: ");    scanf("%d", &n);
```

```
    printf("Enter the page reference string: ");    for (int i = 0; i < n; i++) {  
scanf("%d", &pages[i]);  
    }
```

```
    printf("Enter the number of frames: ");    scanf("%d", &frames);
```

```
    printf("\nFIFO Page Replacement:\n");    fifo(pages, n, frames);
```

```

printf("\nLRU Page Replacement:\n");    lru(pages, n, frames);
return 0;
}

void fifo(int pages[], int n, int frames) {    int frame[MAX_FRAMES];
    int page_faults = 0;    int current_frame = 0;

    for (int i = 0; i < frames; i++) {        frame[i] = -1; // Initialize frames as
empty
    }

    for (int i = 0; i < n; i++) {        int page = pages[i];        int found = 0;

        // Check if page is already in a frame        for (int j = 0; j < frames;
j++) {            if (frame[j] == page) {                found = 1;                break;
            }
        }

        // Page fault: replace the oldest page
        if (!found) {            printf("Page %d caused a page fault.\n", page);
frame[current_frame] = page;            current_frame = (current_frame +
1) % frames; // Circular queue
            page_faults++;
        }

        // Display current state of frames        printf("Frames: ");        for
(int j = 0; j < frames; j++) {            printf("%d ", frame[j]);
        }        printf("\n");
    }

    printf("Total page faults: %d\n", page_faults);
}

void lru(int pages[], int n, int frames) {    int frame[MAX_FRAMES];
    int page_faults = 0;    int counter[MAX_FRAMES] = {0};

    for (int i = 0; i < frames; i++) {

```

```

    frame[i] = -1; // Initialize frames as empty
}

for (int i = 0; i < n; i++) {    int page = pages[i];    int found = 0;

    // Check if page is already in a frame    for (int j = 0; j < frames;
j++) {        if (frame[j] == page) {            found = 1;
counter[j] = i + 1; // Update counter to indicate recent access
            break;
        }
    }

    // Page fault: replace the least recently used page    if (!found) {
printf("Page %d caused a page fault.\n", page);        int min_counter =
counter[0];        int min_index = 0;        for (int j = 1; j < frames; j++)
{            if (counter[j] < min_counter) {                min_counter =
counter[j];
                min_index = j;
            }
        }
        frame[min_index] = page;        counter[min_index] = i + 1; //
Update counter to indicate recent access        page_faults++;
    }

    // Display current state of frames    printf("Frames: ");    for
(int j = 0; j < frames; j++) {        printf("%d ", frame[j]);
    }    printf("\n");
}

printf("Total page faults: %d\n", page_faults);
}

```

Output:

```
Enter the number of pages: 10
Enter the page reference string: 1 2 3 4 1 2 5 1 2 3
Enter the number of frames: 3
```

```
FIFO Page Replacement:
```

```
Page 1 caused a page fault.
```

```
Frames: 1 -1 -1
```

```
Page 2 caused a page fault.
```

```
Frames: 1 2 -1
```

```
Page 3 caused a page fault.
```

```
Frames: 1 2 3
```

```
Page 4 caused a page fault.
```

```
Frames: 4 2 3
```

```
Page 1 caused a page fault.
```

```
Frames: 4 1 3
```

```
Page 2 caused a page fault.
```

```
Frames: 4 1 2
```

```
Page 5 caused a page fault.
```

```
Frames: 5 1 2
```

```
Frames: 5 1 2
```

```
Frames: 5 1 2
```

```
Page 3 caused a page fault.
```

```
Frames: 5 3 2
```

```
Total page faults: 8
```

```
LRU Page Replacement:
Page 1 caused a page fault.
Frames: 1 -1 -1
Page 2 caused a page fault.
Frames: 1 2 -1
Page 3 caused a page fault.
Frames: 1 2 3
Page 4 caused a page fault.
Frames: 4 2 3
Page 1 caused a page fault.
Frames: 4 1 3
Page 2 caused a page fault.
Frames: 4 1 2
Page 5 caused a page fault.
Frames: 5 1 2
Frames: 5 1 2
Frames: 5 1 2
Page 3 caused a page fault.
Frames: 3 1 2
Total page faults: 8
```

Conclusion:

In summary, page replacement policies such as FIFO and LRU play a crucial role in virtual memory management by determining which pages to evict from memory when page faults occur. While FIFO is simple to implement and has low overhead, it may perform poorly in scenarios with irregular access patterns and can suffer from Belady's anomaly. LRU, on the other hand, provides better performance by considering the actual access history of pages, but it comes with higher implementation complexity and overhead. The choice between FIFO and LRU depends on factors such as the system's requirements, workload characteristics, and available resources.

Experiment no: 10.

Name: Meet Brijwani

Roll no: 14

Batch: S11

Aim: Write a program in C to do disk scheduling - FCFS, SCAN, C-SCAN

Theory:

File management and I/O (Input/Output) management are crucial components of operating systems responsible for handling the storage and retrieval of data from disk storage devices efficiently. Disk scheduling algorithms play a vital role in optimizing I/O operations by determining the order in which disk requests are serviced. Among the various disk scheduling algorithms, First-Come, First-Served (FCFS), SCAN, and C-SCAN are widely used. Let's delve into these concepts and understand how each algorithm works and its implications in disk management.

File Management:

File management involves organizing and managing files on disk storage devices to facilitate efficient storage, retrieval, and manipulation of data. It encompasses various operations such as file creation, deletion, access control, and directory management. A file system is responsible for managing these operations and maintaining the structure and integrity of files and directories.

I/O Management:

I/O management deals with managing Input/Output operations between the CPU, memory, and I/O devices such as disks, printers, and network interfaces. Disk I/O operations are particularly significant as they involve relatively slow mechanical devices compared to the CPU and memory. Disk scheduling algorithms are employed to optimize the order in which

disk requests are serviced to minimize seek time and maximize disk throughput.

Disk Scheduling Algorithms:

1. First-Come, First-Served (FCFS):

FCFS is the simplest disk scheduling algorithm, where disk requests are serviced in the order they arrive. It operates on the principle of fairness, as requests are processed based on their arrival times without any consideration for their locations on the disk.

Mechanism:

When a disk request arrives, it is added to the end of the request queue.

The disk scheduler services requests in the order they are queued.

Each request is processed sequentially, starting from the innermost track to the outermost track of the disk.

Implications:

FCFS is easy to implement and ensures fairness in servicing requests.

However, it may lead to longer seek times, especially if requests are scattered across the disk, resulting in poor disk performance.

2. SCAN (Elevator) Algorithm:

SCAN, also known as the elevator algorithm, simulates the movement of an elevator moving up and down a building. It services requests in one direction until reaching the end of the disk, then reverses direction and continues servicing requests in the opposite direction.

Mechanism:

The disk head starts from one end of the disk and moves towards the other end while servicing requests along its path.

When it reaches the end of the disk, it reverses direction and starts moving towards the opposite end, servicing requests along the way.

SCAN prevents the disk head from unnecessarily traversing the entire disk by changing direction at the disk boundaries.

Implications:

SCAN reduces the average seek time by prioritizing requests closer to the current position of the disk head.

However, it may result in starvation for requests located at the extremes of the disk if there is a continuous stream of requests in one direction.

3. C-SCAN (Circular SCAN) Algorithm:

C-SCAN is an enhancement of the SCAN algorithm that overcomes the potential starvation issue by treating the disk as a circular buffer. After reaching one end of the disk, the disk head jumps to the other end without servicing requests, ensuring fairness in request servicing.

Mechanism:

Similar to SCAN, the disk head moves in one direction, servicing requests until reaching the end of the disk.

Instead of reversing direction immediately, C-SCAN jumps to the other end of the disk without servicing requests.

It then continues servicing requests in the same direction, preventing starvation for requests located at the disk boundaries.

Implications:

C-SCAN provides fairness in request servicing by preventing starvation for requests at the extremes of the disk.

However, it may result in slightly higher average seek times compared to SCAN due to the jump between disk ends.

Program:

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#define MAX_REQUESTS 100 // Maximum number of disk requests
```

```
#define MAX_CYLINDERS 200 // Maximum number of cylinders on the disk
```

```
// Function prototypes void fcfs(int requests[], int n, int initial_position);  
void scan(int requests[], int n, int initial_position, int cylinders); void  
c_scan(int requests[], int n, int initial_position, int cylinders);
```

```
int main() {    int requests[MAX_REQUESTS];  
    int n, initial_position, cylinders;
```

```
    printf("Enter the number of disk requests: ");    scanf("%d", &n);
```

```
    printf("Enter the disk requests: ");    for (int i = 0; i < n; i++) {  
scanf("%d", &requests[i]);  
    }
```

```
    printf("Enter the initial position of the disk head: ");    scanf("%d",  
&initial_position);
```

```
    printf("Enter the total number of cylinders on the disk: ");  
scanf("%d", &cylinders);
```

```
    printf("\nFirst-Come, First-Served (FCFS):\n");    fdfs(requests, n,  
initial_position);
```

```
    printf("\nSCAN:\n");    scan(requests, n, initial_position, cylinders);
```

```
    printf("\nC-SCAN:\n");    c_scan(requests, n, initial_position,  
cylinders);
```

```
    return 0;  
}
```

```
void fdfs(int requests[], int n, int initial_position) {    int total_seek_time =  
0;
```

```
    printf("Sequence of servicing requests:\n");    for (int i = 0; i < n; i++) {  
printf("%d ", requests[i]);    total_seek_time += abs(requests[i] -  
initial_position);    initial_position = requests[i];  
    }
```

```
    printf("\nTotal seek time: %d\n", total_seek_time);  
}
```

```
void scan(int requests[], int n, int initial_position, int cylinders) {    int  
total_seek_time = 0;    int direction = 1; // 1 for right, -1 for left
```

```
    printf("Sequence of servicing requests:\n");
```

```

    // Sort requests to service in ascending order    for (int i = 0; i < n - 1;
i++) {        for (int j = 0; j < n - i - 1; j++) {            if (requests[j] >
requests[j + 1]) {                int temp = requests[j];                requests[j] =
requests[j + 1];                requests[j + 1] = temp;
            }
        }
    }

```

```

    // Find the index where the disk head should reverse direction    int
reverse_index = 0;    while (reverse_index < n &&
requests[reverse_index] < initial_position) {
        reverse_index++;
    }

```

```

    // Service requests in one direction    for (int i = reverse_index - 1; i
>= 0; i--) {        printf("%d ", requests[i]);        total_seek_time +=
abs(requests[i] - initial_position);        initial_position = requests[i];
    }

```

```

    // Service requests in the reverse direction    for (int i = reverse_index;
i < n; i++) {        printf("%d ", requests[i]);        total_seek_time +=
abs(requests[i] - initial_position);        initial_position = requests[i];
    }

```

```

    printf("\nTotal seek time: %d\n", total_seek_time);
}

```

```

void c_scan(int requests[], int n, int initial_position, int cylinders) {    int
total_seek_time = 0;

```

```

    printf("Sequence of servicing requests:\n");

```

```

    // Sort requests to service in ascending order    for (int i = 0; i < n - 1;
i++) {        for (int j = 0; j < n - i - 1; j++) {            if (requests[j] >
requests[j + 1]) {                int temp = requests[j];                requests[j] =
requests[j + 1];                requests[j + 1] = temp;
            }
        }
    }

```

```

    }
}

// Service requests in one direction    for (int i = 0; i < n && requests[i]
< initial_position; i++) {        printf("%d ", requests[i]);
total_seek_time += abs(requests[i] - initial_position);        initial_position
= requests[i];
    }

// Jump to the beginning of the disk    printf("%d ", 0);
total_seek_time += initial_position;

// Service requests in the reverse direction    for (int i = n - 1; i >= 0
&& requests[i] >= initial_position; i--) {        printf("%d ", requests[i]);
        total_seek_time += abs(requests[i] - initial_position);
initial_position = requests[i];
    }

printf("\nTotal seek time: %d\n", total_seek_time);
}

```

Output:

```
Enter the number of disk requests: 5
Enter the disk requests: 98 183 37 122 14
Enter the initial position of the disk head: 53
Enter the total number of cylinders on the disk: 200

First-Come, First-Served (FCFS):
Sequence of servicing requests:
98 183 37 122 14
Total seek time: 469

SCAN:
Sequence of servicing requests:
37 14 98 122 183
Total seek time: 208

C-SCAN:
Sequence of servicing requests:
14 0 183
Total seek time: 222

...Program finished with exit code 0
Press ENTER to exit console.
```

Conclusion:

File management and I/O management are essential components of operating systems, responsible for efficient storage and retrieval of data from disk storage devices. Disk scheduling algorithms such as FCFS, SCAN, and C-SCAN play a crucial role in optimizing disk I/O operations by determining the order in which disk requests are serviced. While FCFS provides simplicity and fairness, SCAN and CSCAN aim to minimize seek times and prevent starvation for requests located at the extremes of the disk. The choice of disk scheduling algorithm depends on factors such as disk workload characteristics, system requirements, and performance considerations.